

Essay summarization and automated grading system

By,

Sai Charan Merugu

Sai Gottumukkala

Swetha Guntupalli

Karishma Bollineni

Problem Statement

Nowadays, instructors are often overwhelmed by the amounts of essay correction they need to grade everyday. Traditional methods of evaluation consumes ample amounts of time and heavily relies on manual efforts. Besides, it can also be inconsistent and biased at times.

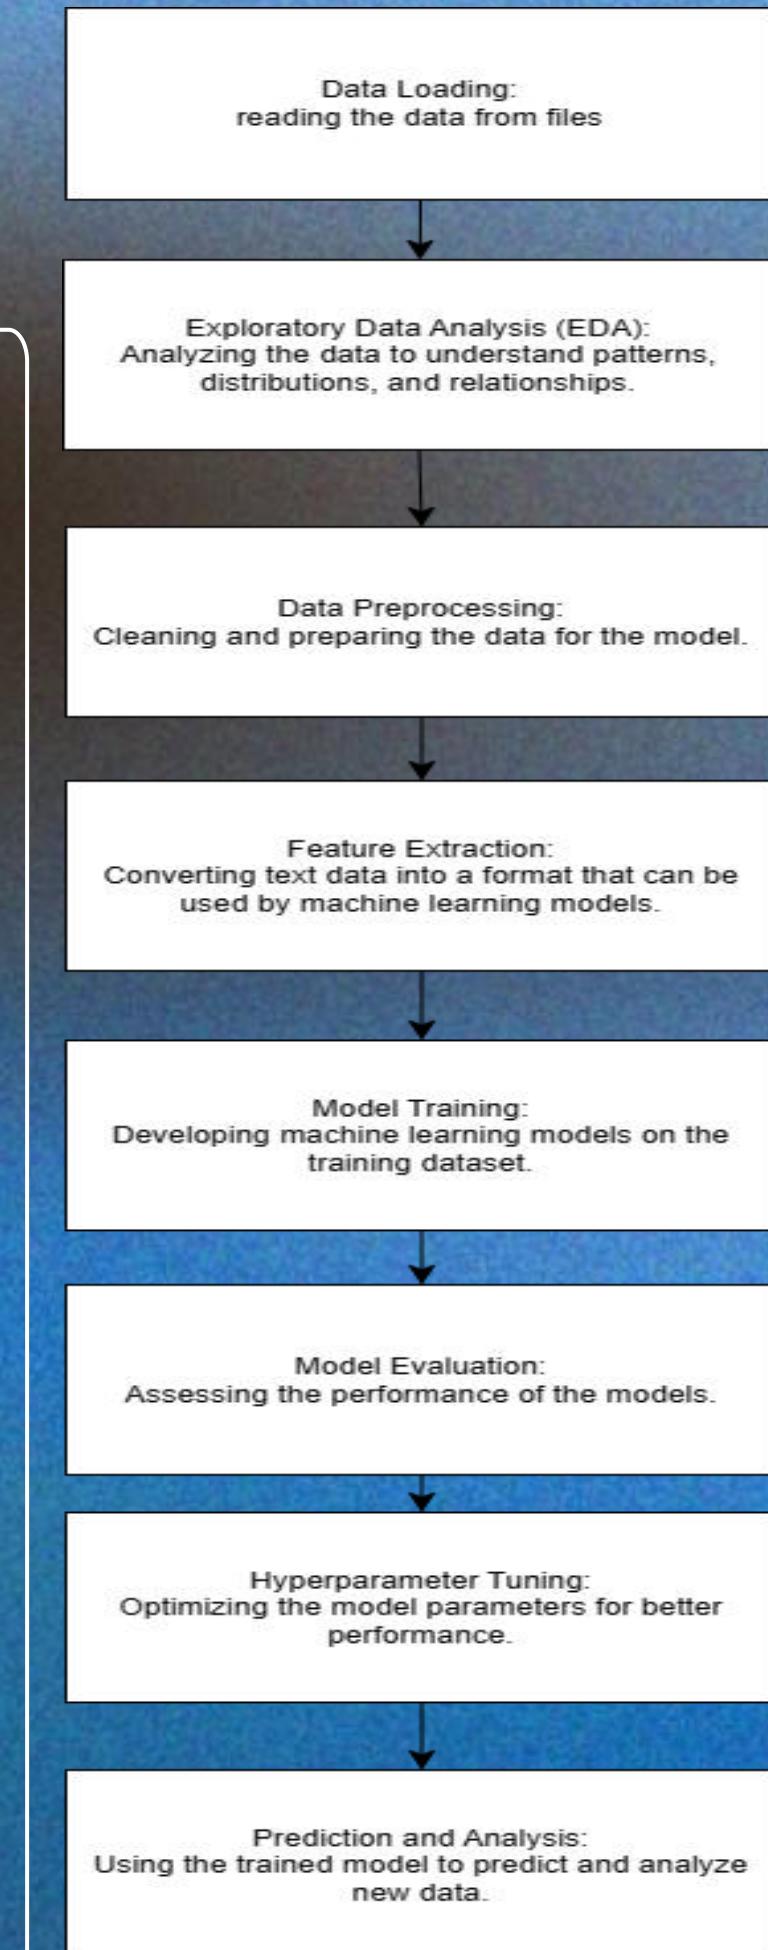
Lack of an efficient system that can do this job instead of the actual person, thereby saving some time, triggered us with this idea.

This paved ways for our project idea that can summarize the essay and grade it accordingly. The summarized essay will serve as a quick review option if the instructor chooses to re-verify the grade assigned.

Model/Methodology

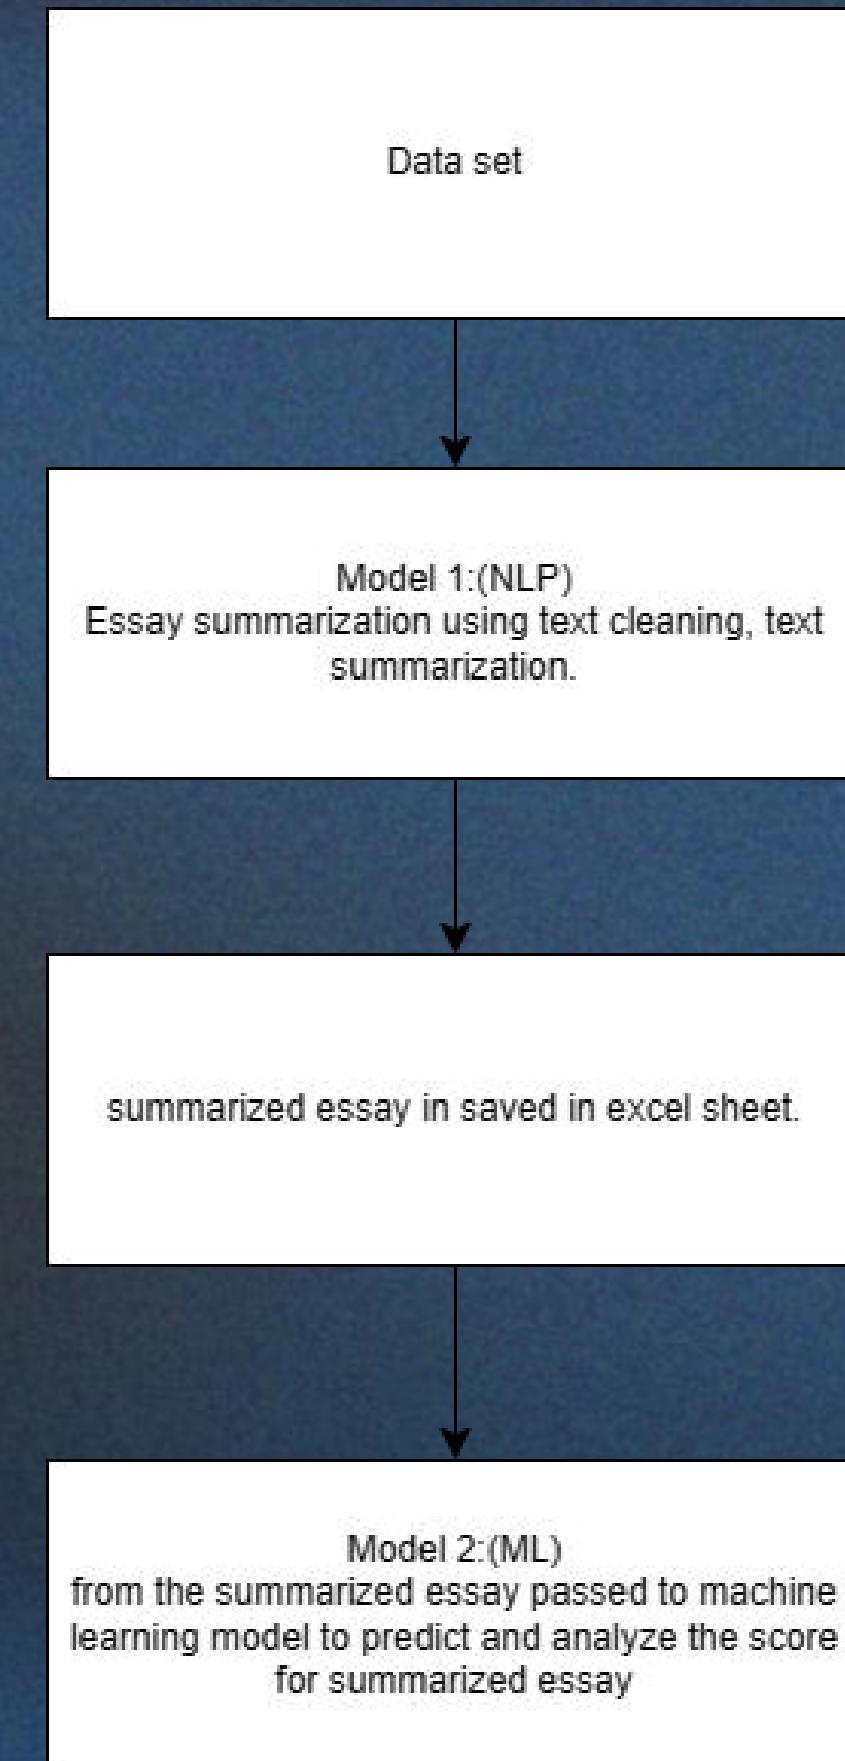
Overall Workflow Diagram

- Data Loading: Reading the data from files.
- Exploratory Data Analysis (EDA): Analyzing the data to understand patterns, distributions, and relationships.
- Data Preprocessing: Cleaning and preparing the data for the model.
- Feature Extraction: Converting text data into a format that can be used by machine learning models.
- Model Training: Developing machine learning models on the training dataset.
- Model Evaluation: Assessing the performance of the models.
- Hyperparameter Tuning: Optimizing the model parameters for better performance.
- Prediction and Analysis: Using the trained model to predict and analyze new data.



Architecture diagram of model

- Model 1:NLP
 - After analyzing the essay ,it undergoes to text cleaning and text summarization. This summarized essay is saved as excel sheet along with score.
- Model 2:Machine Learning:
 - The summarized essay saved from excel under go preprocessing ,then passed to model . The model will learning from the data and predict the score based on the summarized essay.



Dataset

The dataset consist of 8 distinct sets of essays, corresponding to different prompts. each essay vary in length from approximately 150 to 550 words and Some essays are based on provided sources, while others are independent. these responses were crafted by students in grades 7 to 10. Each essay has been hand-graded and has received scores from multiple raters for increased reliability.

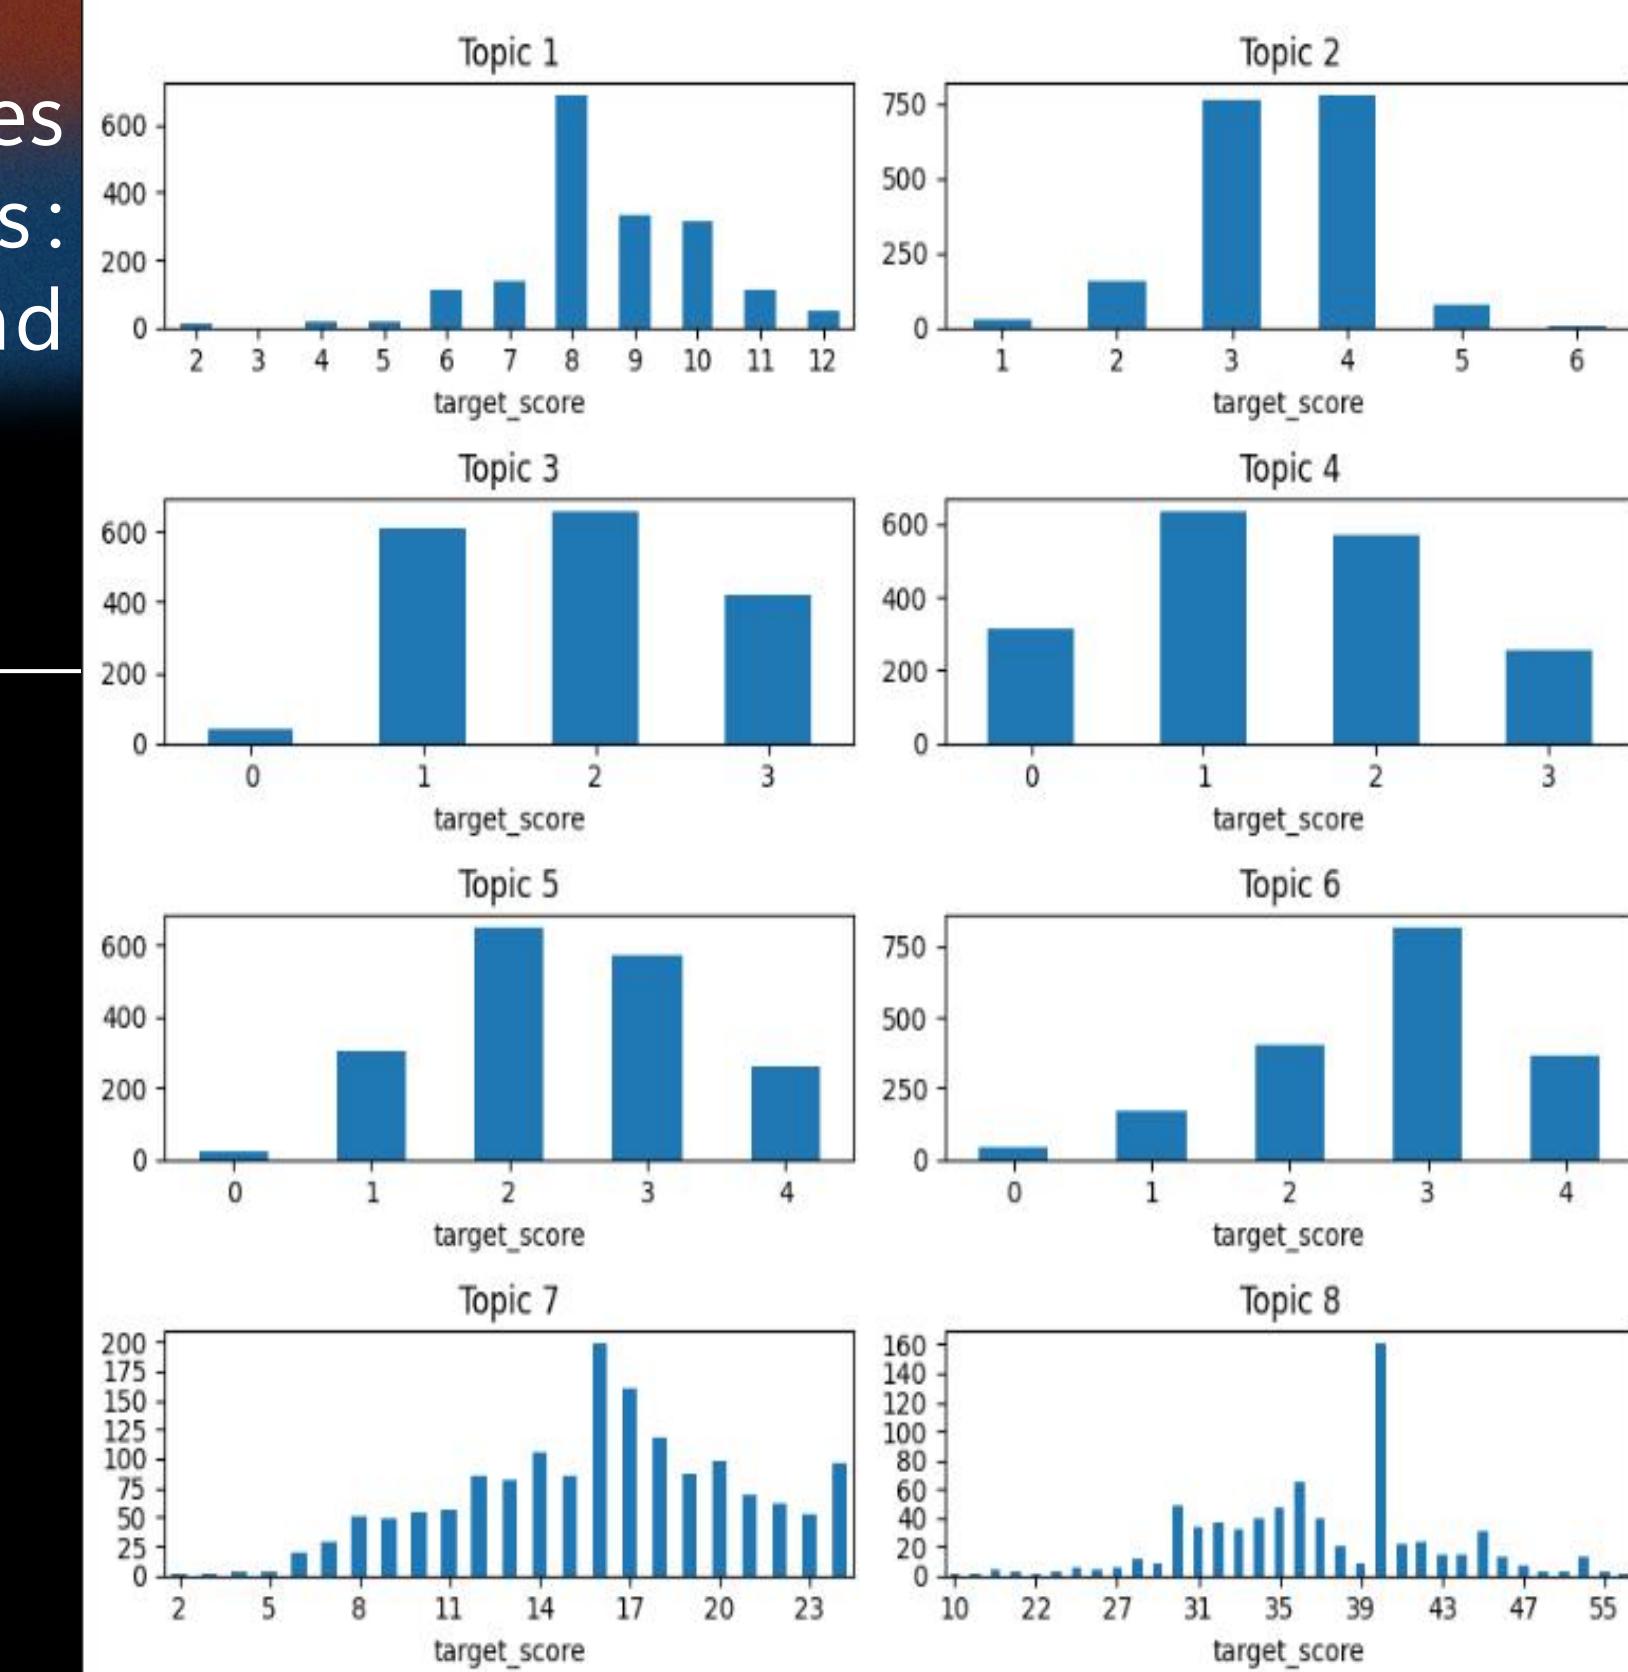
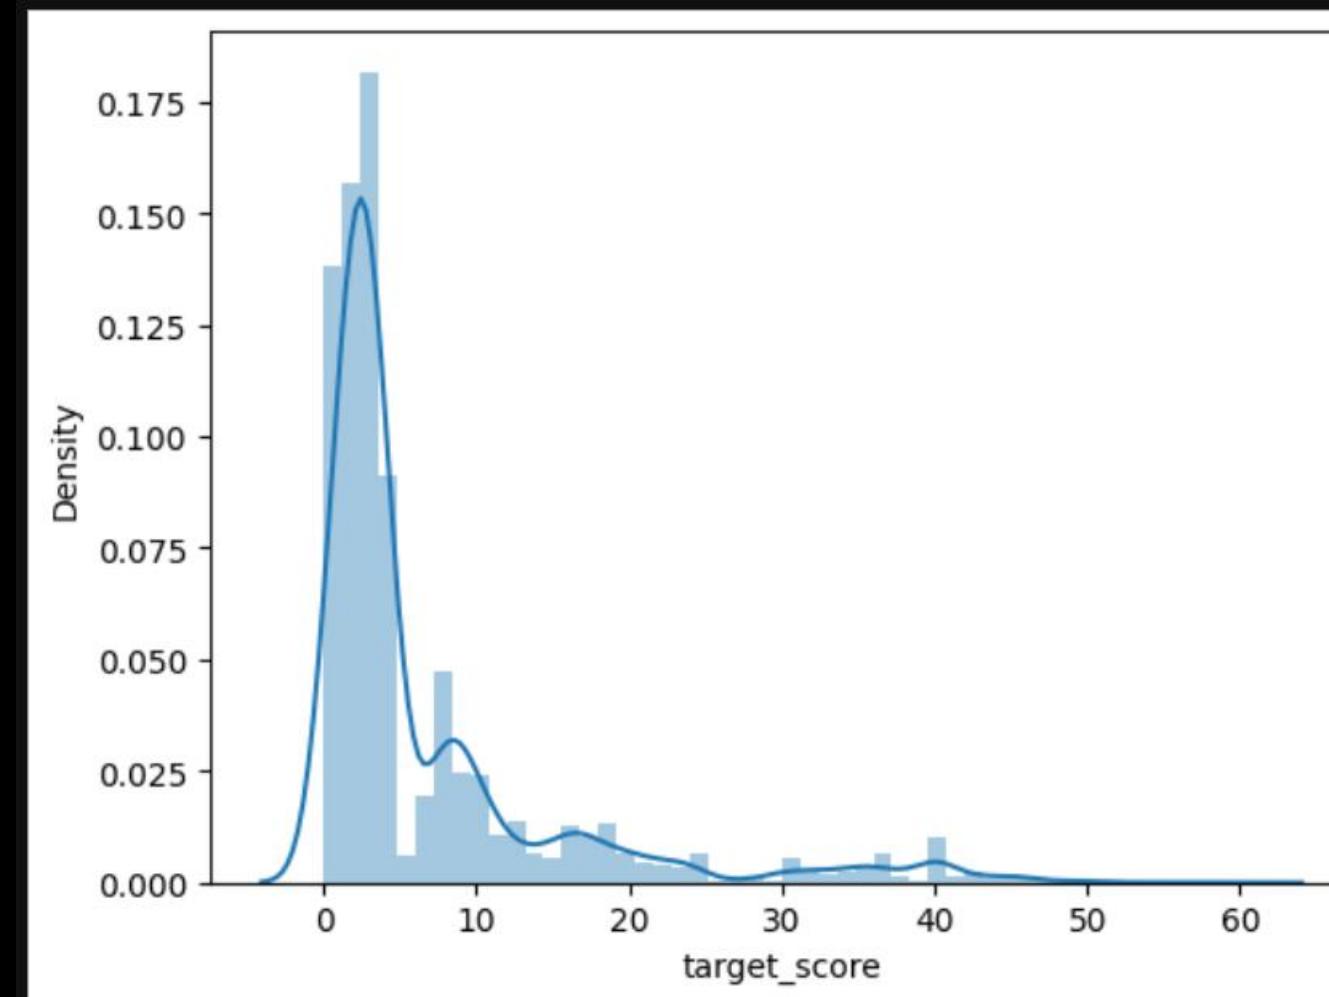
The dataset has following columns:

- Essay_id: A unique identifier for each student essay.
- essay_set: A numerical identifier (1-8) indicating the set to which an essay belongs.
- essay: The ASCII text content of the student's response.
- rater1_domain1: The first rater's score for the primary domain or aspect of the essay.
- rater2_domain1: The second rater's score for the primary domain.
- rater3_domain1: The third rater's score for the primary domain, applicable only for some essays in set 8.
- domain1_score: The final agreed-upon score for the primary domain after resolving differences between raters.
- rater1_domain2, rater2_domain2: Scores from two raters for a secondary domain, applicable only to essays in set 2.
- domain2_score: The resolved score for the secondary domain, only present for essays in set 2.
- Trait Scores (rater1_trait1 to rater3_trait6): Specific characteristic scores for essays in sets 7-8. Each trait score evaluates different aspects of writing, such as organization, vocabulary, grammar, etc

Design of Features/labels with diagram

By comparing the features and concluded useful features from the dataset, categorized the features: essay_id,topic,essay,target_score of the data set ,and plotted the 8 essay topics by target score.

Analyzed the score of the total essays, plot the graph.



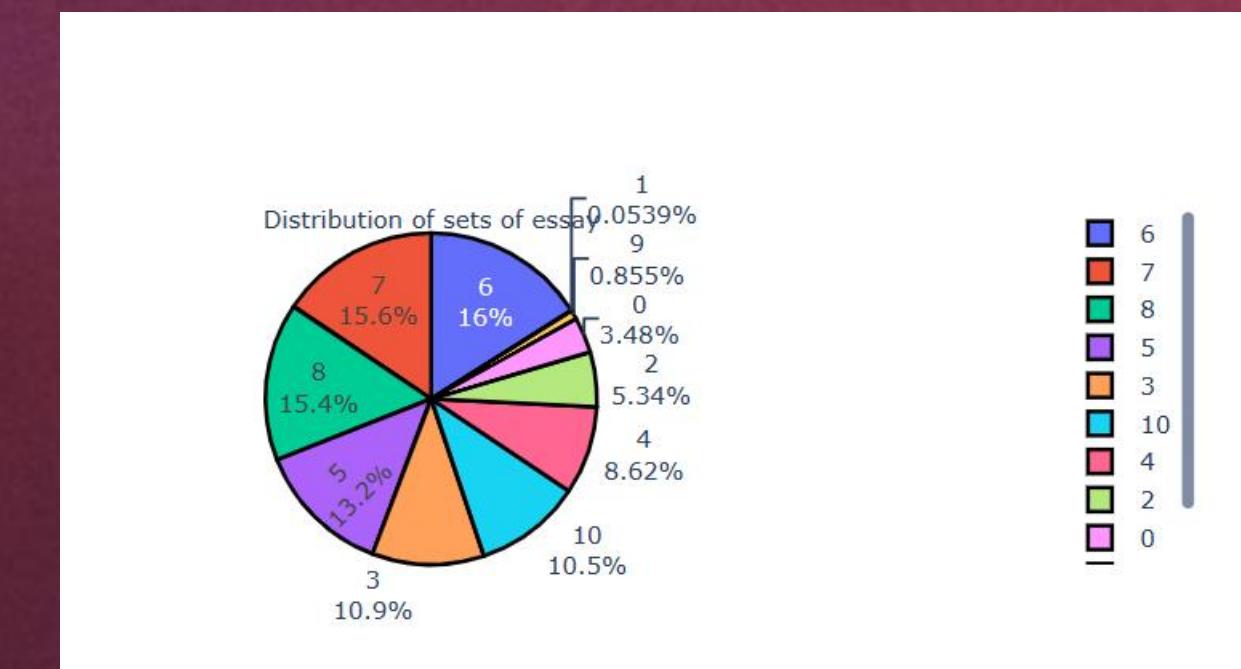
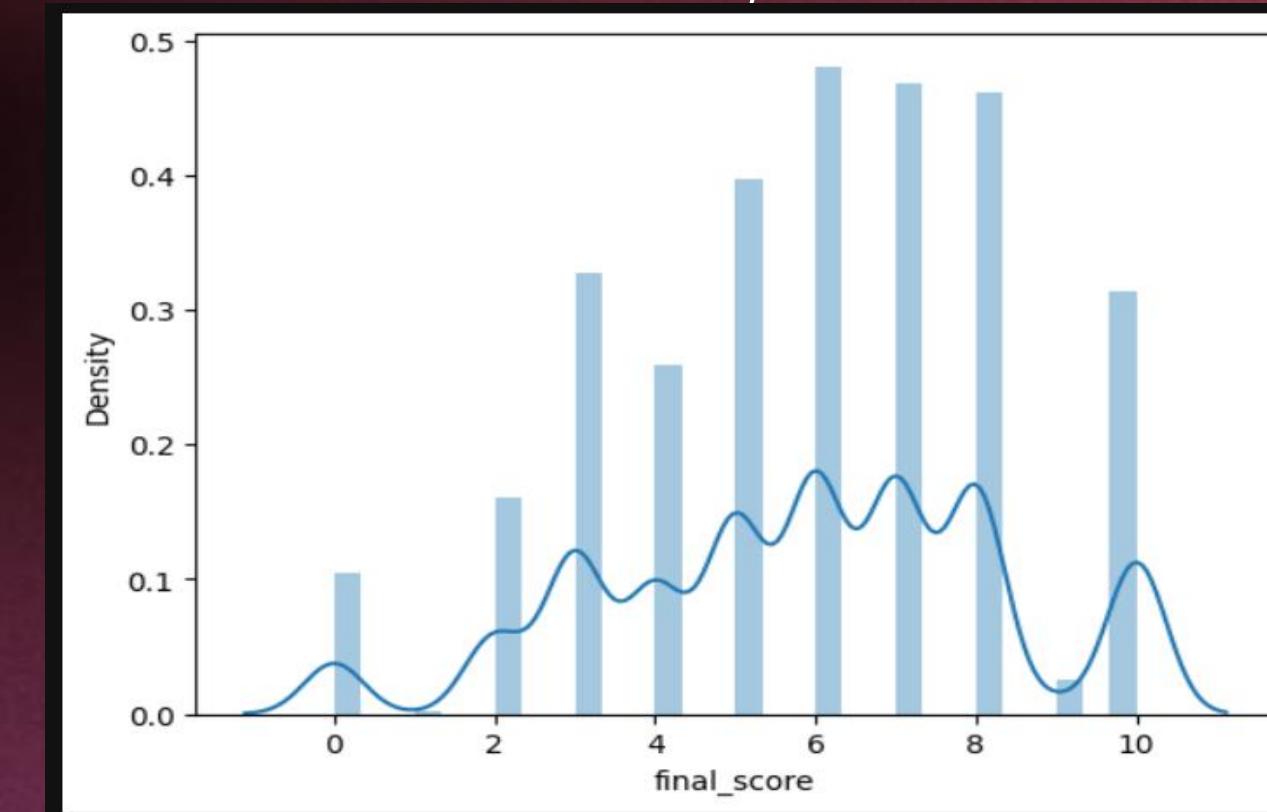
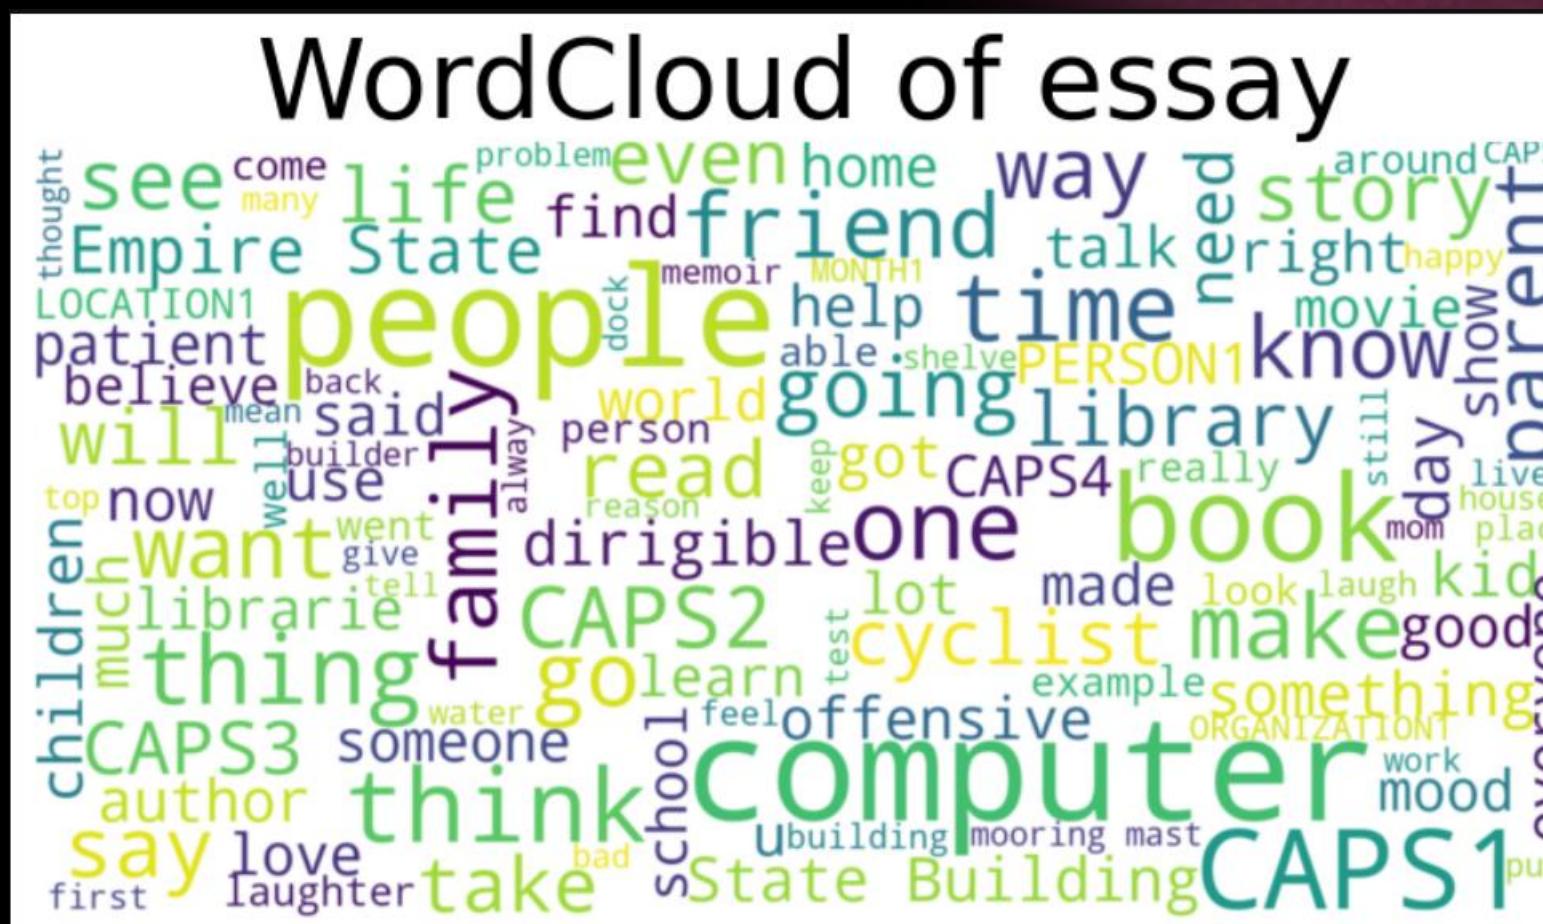
Data Visualization and data preprocessing

Data preprocessing: as the target score is high in some columns and low in some columns, we normalized the and stored in final_score.

```
: min_range = [2,1,0,0,0,0,0,0]
: max_range = [12,6,3,3,4,4,30,60]

def normalize(x,mi,ma):
    #print("Before Normalization: "+str(x))
    x = (x-mi)/(ma-mi)
    #print("After Normalization : "+str(x))
    return round(x*10)

df['final_score']=df.apply(lambda x:normalize(x['target_score'],min_range[x['topic']-1],max_range[x['topic']-1]),axis=1)
```



Data Visualization and data preprocessing

Data preprocessing: data cleaning for the given essay using regular expression clean html tags, special symbols, slashes, on word punctuation ,numbers, white spaces and convert them to lower case.

For the clean text in the essay ,we process the text summarization from the learning concept of the course using spacy library , using stop words and punctuation cleaned further text and stored in summary field and exported the excel file.

```
import re

def clean_text(text):
    # Remove placeholders like @CAPS1
    text = re.sub(r'@\w+', '', text)
    # Remove HTML tags if present
    text = re.sub(r'<.*?>', '', text)
    # Handle specific cases like 'skills/affects'
    # Option 1: Replace slashes with spaces to separate words
    text = re.sub(r'/', ' ', text)
    # Option 2: If you prefer to keep words joined post-slash removal, uncomment the following line
    text = re.sub(r'/', '', text)
    # Replace non-word characters (punctuation) with spaces except for apostrophes
    text = re.sub(r'[\^\\w\\s\\']', ' ', text)
    # Remove numbers if they are not relevant
    text = re.sub(r'\d+', '', text)
    # Remove any underscores - often used as placeholders
    text = re.sub(r'_', ' ', text)
    # Normalize white space to a single space
    text = re.sub(r'\\s+', ' ', text)
    # Convert to lowercase to maintain uniformity
    text = text.lower().strip()
return text
```

```
import spacy
from spacy.lang.en.stop_words import STOP_WORDS
from string import punctuation
import pandas as pd
from heapq import nlargest

# Initialize spacy
nlp = spacy.load('en_core_web_sm')
stopwords = list(STOP_WORDS)
punctuation += '\n' # Append newline to punctuation

def summarize_essay(essay, select_length=0.3):
    """ Summarize the essay using spacy for NLP processing. """
    doc = nlp(essay)
    word_frequencies = {}

    # Calculate word frequencies, ignoring stopwords and punctuation
    for word in doc:
        if word.text.lower() not in stopwords and word.text.lower() not in punctuation:
            word_frequencies[word.text] = word_frequencies.get(word.text, 0) + 1

    # Normalize frequencies and compute sentence scores
    max_frequency = max(word_frequencies.values(), default=1)
    sentence_scores = {}
    for sent in doc.sents:
        for word in sent:
            if word.text.lower() in word_frequencies:
                sentence_scores[sent] = sentence_scores.get(sent, 0) + word_frequencies[word.text.lower()]

    # Select the top sentences based on their scores
    summary_sentences = nlargest(int(len(list(doc.sents)) * select_length), sentence_scores, key=sentence_scores.get)
    summary = ' '.join(sent.text for sent in summary_sentences)

    return summary

def process_and_save_summaries(input_file, output_file_path):
    """ Read essays from an Excel file, summarize them, and save the summaries to a new Excel file. """
    essay_data = input_file
    essay_data['summary'] = essay_data['cleaned'].apply(summarize_essay)

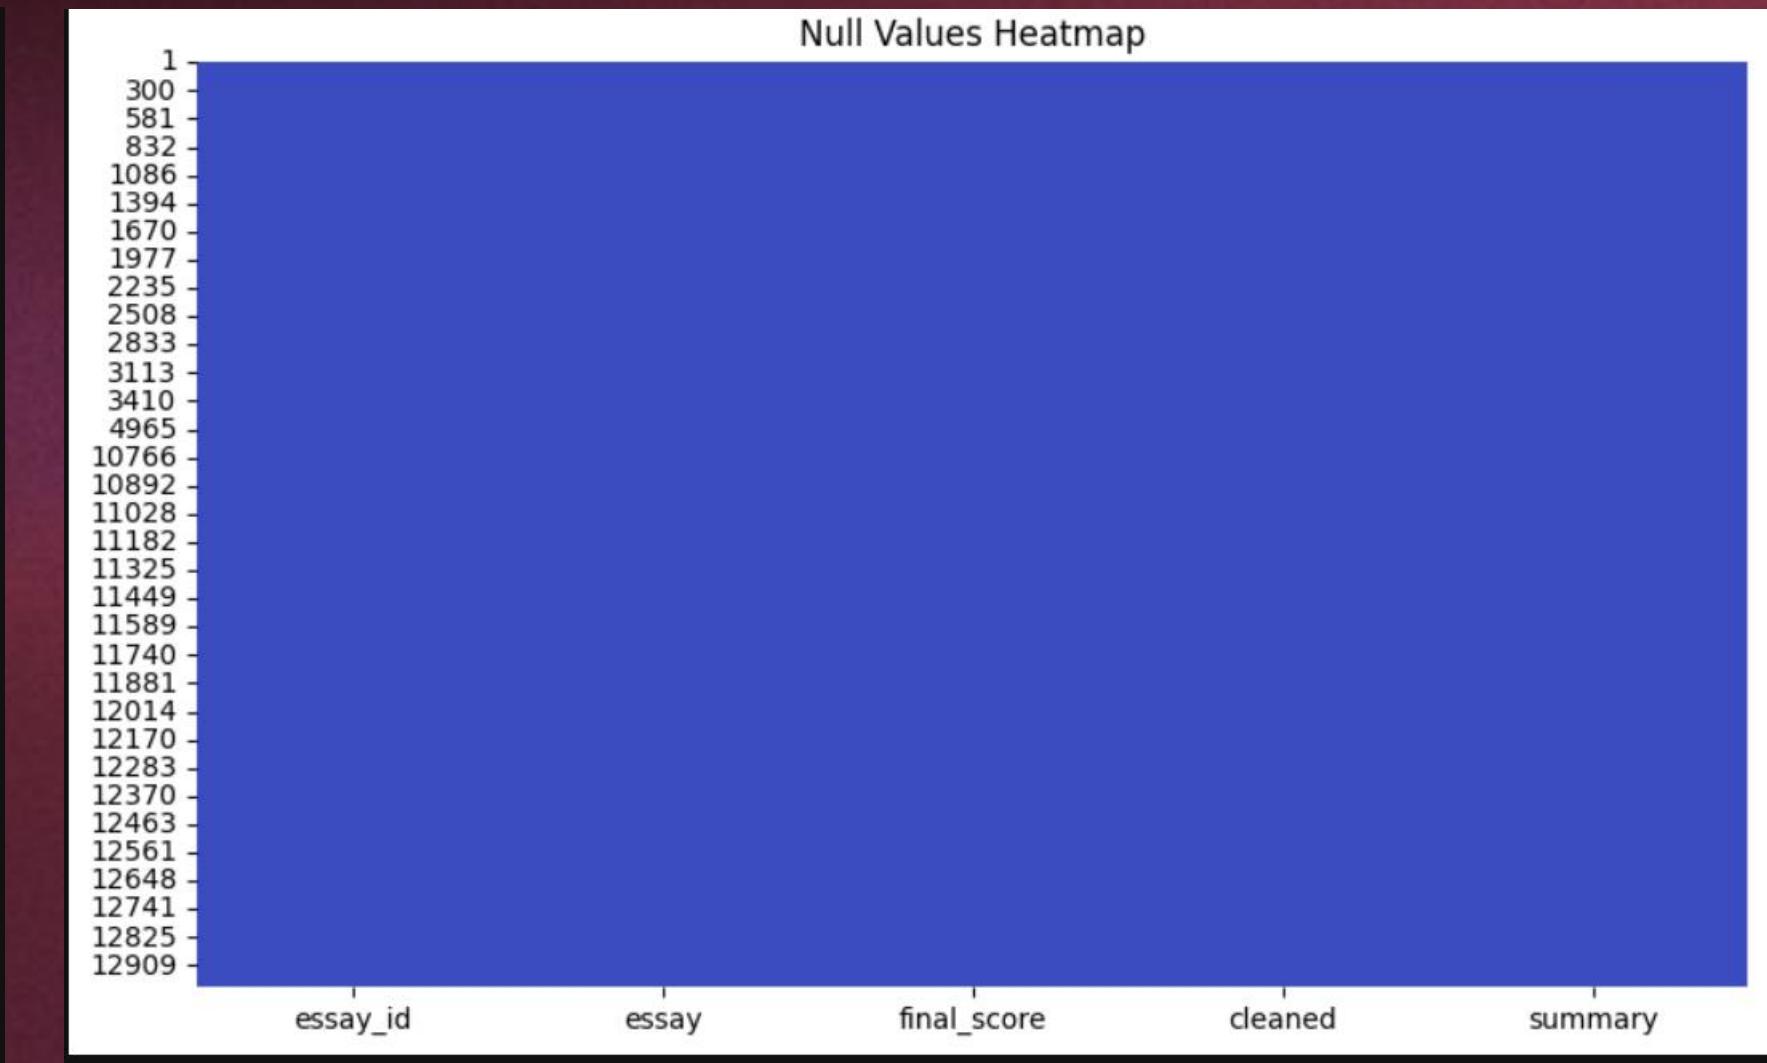
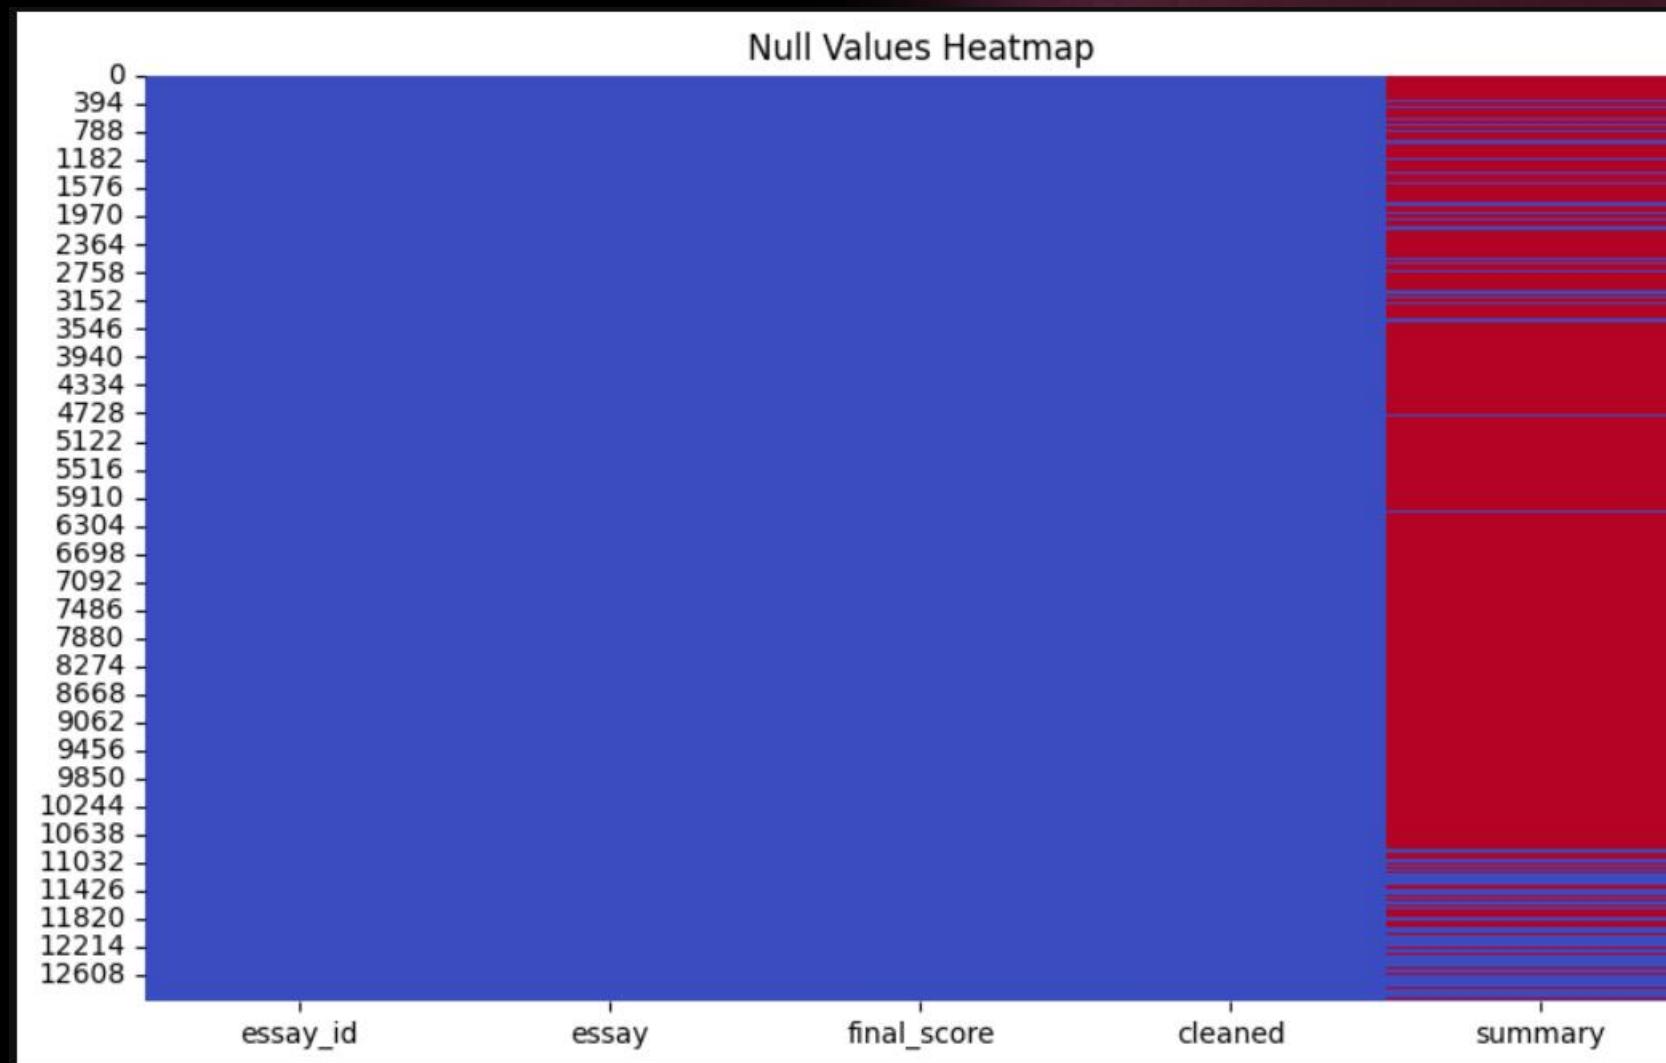
    # Save the updated DataFrame to a new Excel file
    essay_data.to_excel(output_file_path, index=False)
    print(f"Summaries saved to {output_file_path}")

    # Specify file paths
    input_file_path = df # Path to the input Excel file with essays
    output_file_path = r'C:\Users\merug\Desktop\nlp_project\summarized_essays.xlsx' # Path for the output Excel file with summaries

    # Process essays and save summaries
    process_and_save_summaries(input_file_path, output_file_path)
```

Data Visualization and data preprocessing

Data preprocessing: after summarizing the essay, before training checking the null fields and cleaning it. Here are the heatmap representation of the analysis.



Implementation

- NLP libraries
- spaCy: It is an open-source NLP library designed for performing complex NLP tasks with ease. spaCy is known for its speed and efficiency and comes with pre-built statistical models and word vectors.
- Regular Expression (re): Regular expressions are used for string searching and manipulation, often for complex patterns of text. For instance, you might use regular expressions to validate email addresses, remove unwanted characters, find all URLs in a text, or perform substitutions in a string.
- String: It's often used in text processing to remove punctuation from text before preprocessing or tokenizing because punctuation marks can interfere with how text is interpreted or analyzed.
- Common Libraries
- Numpy: It's heavily used in numerical computations, data transformation, and is a key library in the ecosystem of scientific computing in Python.
- Pandas: It provides data structures like DataFrames and Series, which are ideal for working with structured data. Pandas is widely used for data processing, cleaning, and analysis.
- Seaborn: It offers a more powerful and simpler interface for creating complex visualizations, including various plot types like scatter plots, barcharts, and violin plots, with sensible default designs.
- matplotlib.pyplot: It's used for creating static, interactive and animated visualizations in Python.

Model 2:Libraries ,algorithm and implementation

Logistic regression:

Logistic Regression is a statistical algorithm used primarily for binary classification problems, which is a type of predictive analysis. Regression models the probabilities for classification problems with two possible outcomes. It's an extension of the linear regression model for classification tasks.

Eg:"Spam" or "NotSpam" in email filtering systems.

"Malignant" or "Benign" for a tumor diagnosis.

We used this algorithm to analyze the summarized essay for classification of scores based on essay. And used metrics for testing model learning: Accuracy, Precision, Recall, F1 Score, confusion matrix

Model 2: Libraries ,algorithm and implementation

Logistic regression:

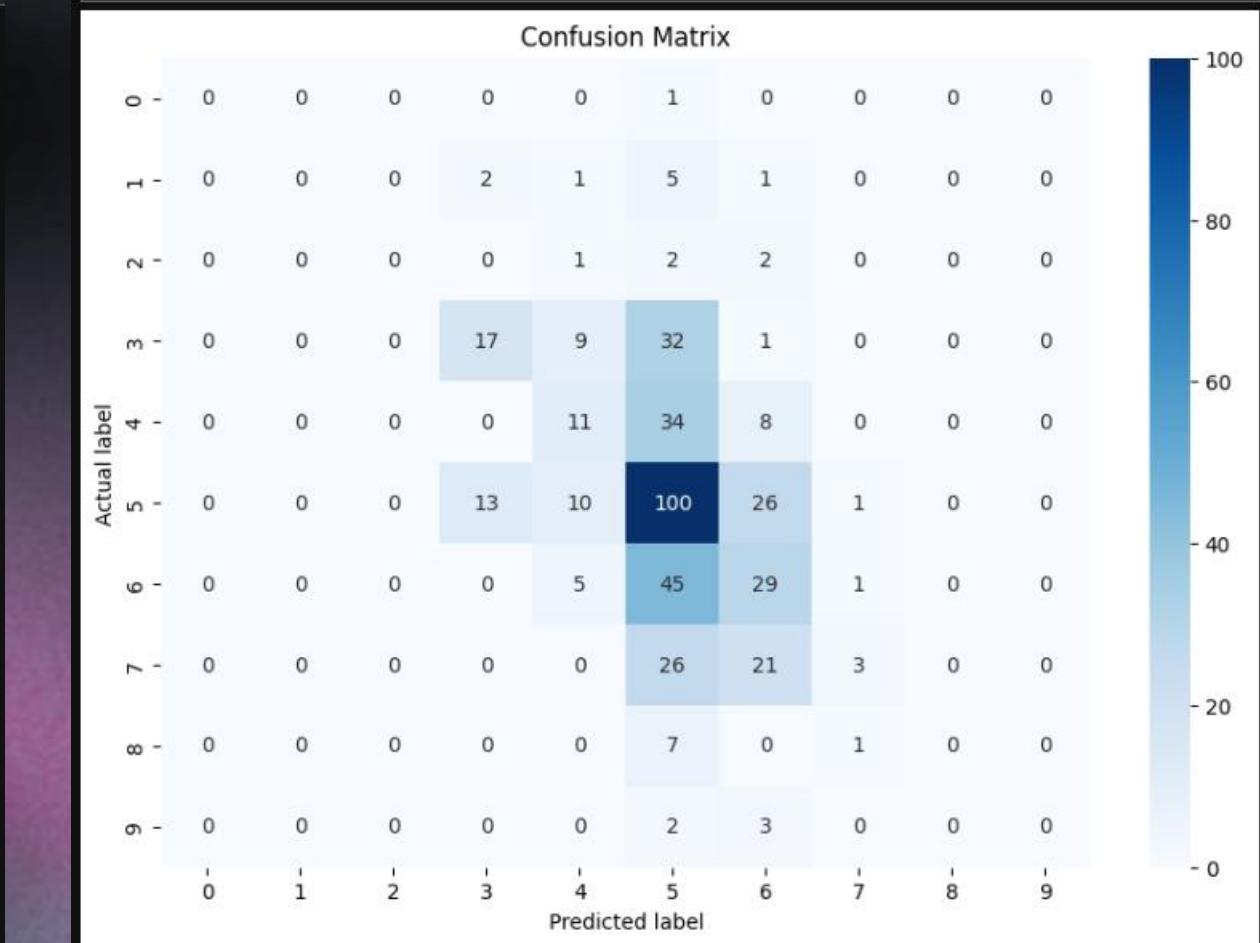
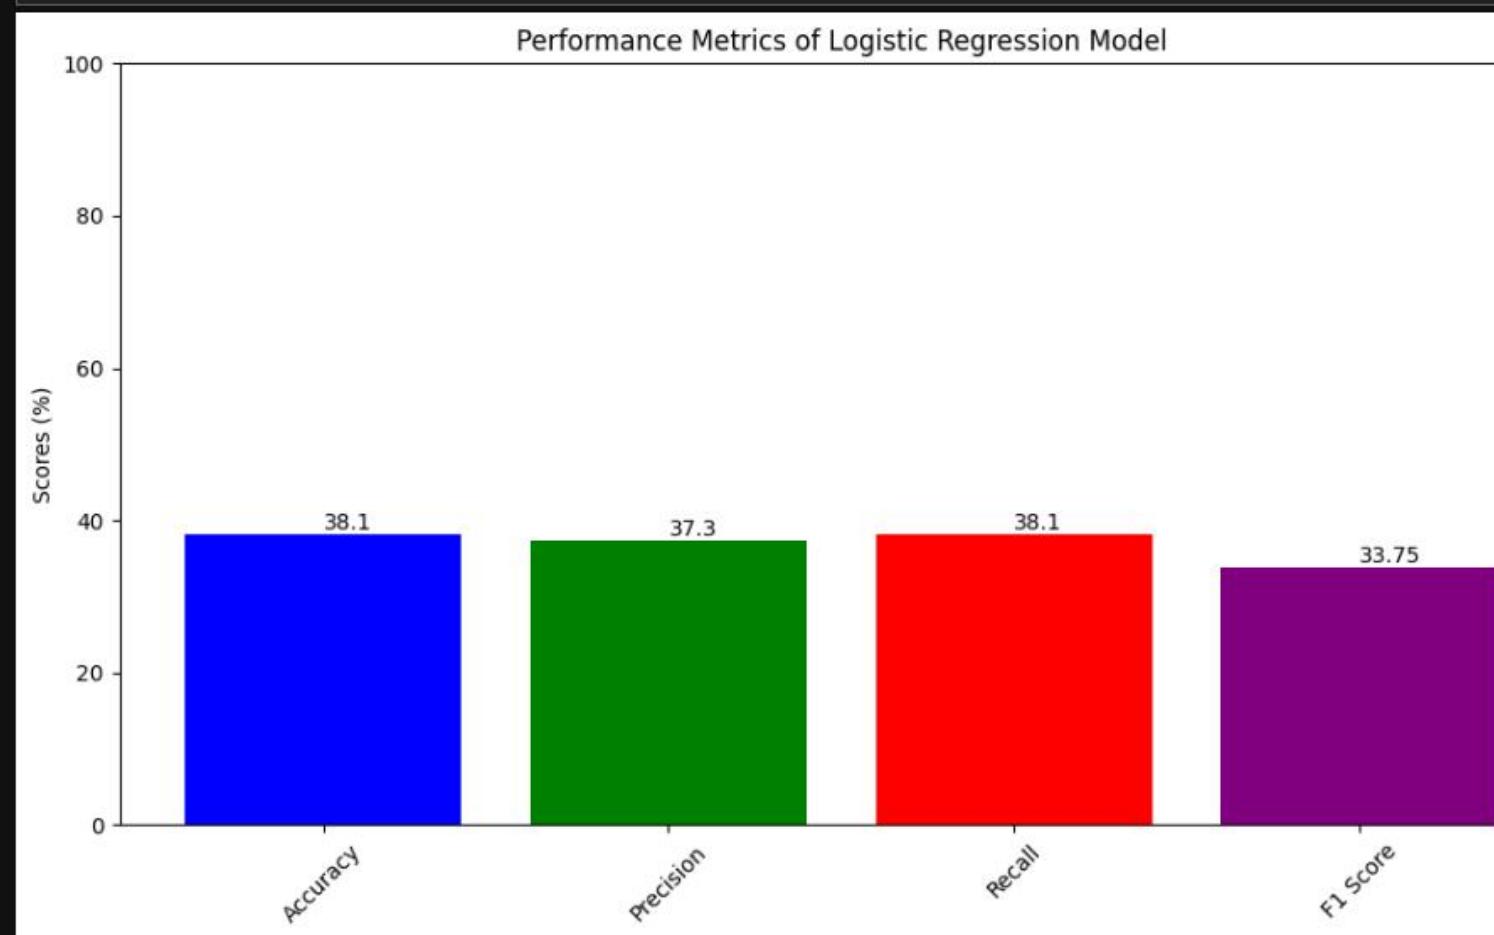
```
: from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report
import pandas as pd

# Split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(df['summary'], df['final_score'], test_size=0.2, shuffle=True, random_state=26)
# Initialize TF-IDF vectorizer
vectorizer = TfidfVectorizer()
# Fit the vectorizer on the training essays
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)
# Initialize and train the Logistic Regression model
logreg_model = LogisticRegression()
logreg_model.fit(X_train_tfidf, y_train)
# Predictions
predictions = logreg_model.predict(X_test_tfidf)
# Evaluation metrics
accuracy = accuracy_score(y_test, predictions)
precision = precision_score(y_test, predictions, average='weighted')
recall = recall_score(y_test, predictions, average='weighted')
f1 = f1_score(y_test, predictions, average='weighted')
confusion = confusion_matrix(y_test, predictions)
classification_rep = classification_report(y_test, predictions)
# Print evaluation metrics
print("Logistic Regression Model (Without Hyperparameter Tuning)")
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("Confusion Matrix:\n", confusion)
print("Classification Report:\n", classification_rep)
```

```
Logistic Regression Model (Without Hyperparameter Tuning)
Accuracy: 0.38095238095238093
Precision: 0.37297650086071066
Recall: 0.38095238095238093
F1 Score: 0.33749662648382217
Confusion Matrix:
[[ 0  0  0  0  0  1  0  0  0  0]
 [ 0  0  0  2  1  5  1  0  0  0]
 [ 0  0  0  0  1  2  2  0  0  0]
 [ 0  0  0  17  9  32  1  0  0  0]
 [ 0  0  0  0  11  34  8  0  0  0]
 [ 0  0  0  13  10  100  26  1  0  0]
 [ 0  0  0  0  5  45  29  1  0  0]
 [ 0  0  0  0  0  26  21  3  0  0]
 [ 0  0  0  0  0  7  0  1  0  0]
 [ 0  0  0  0  0  2  3  0  0  0]]
```

Model 2: Libraries ,algorithm and implementation

Logistic regression:



Model 2: Libraries ,algorithm and implementation

Logistic regression: we divided the dataset in 80 for training and testing.
TfidfVectorizer from sklearn.feature_extraction.text is used to convert the text data into numerical features suitable for model training.

A LogisticRegression model is instantiated and trained on the vectorized training data (X_train_tfidf).
The trained model is used to predict the scores of the vectorized test data (X_test_tfidf).

Accuracy: The proportion of total predictions that were correct.

Precision (Weighted): The ability of the classifier not to label a negative sample a positive, with consideration for the imbalance in class distribution.

Recall (Weighted): The ability of the classifier to find all the positive samples, again weighted for class imbalance.

F1 Score (Weighted): A weighted average of precision and recall, which balances both metrics.

A confusion matrix is generated, showing true positives, false positives, true negatives, and false negatives.

Decision tree

Decision Tree:

Decision Trees are a type of Supervised Machine Learning algorithm that is predominantly used for classification problems, but can also be adapted for regression.

A Decision Tree algorithm segments the dataset into branches to form a tree structure. Each internal node of the tree represents a test on an attribute (e.g., whether a coinflip comes up heads or tails), each branch represents the outcome of the test, and each leaf node represents a class label (decision taken after computing all attributes)

We used this algorithm to analyze the summarized essay for classification of score based on essay. And used metrics for testing model learning: Accuracy, Precision , Recall , F1 Score, confusion matrix

Decision tree

Decision Tree:

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confusion_matrix, classification_report
import pandas as pd

# Load the dataset containing essays and summaries

# Initialize TF-IDF vectorizer
vectorizer = TfidfVectorizer()

# Fit the vectorizer on the training essays
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)

# Initialize and train the Decision Tree model
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train_tfidf, y_train)

# Predictions
predictions = dt_model.predict(X_test_tfidf)

# Evaluation metrics
accuracy = accuracy_score(y_test, predictions)
precision = precision_score(y_test, predictions, average='weighted')
recall = recall_score(y_test, predictions, average='weighted')
f1 = f1_score(y_test, predictions, average='weighted')
# AUC-ROC score
# try:
#     auc_roc = roc_auc_score(y_test, predictions, average='weighted', multi_class='ovr')
# except ValueError:
#     auc_roc = 0.0 # Set to 0 if AUC-ROC score calculation fails

confusion = confusion_matrix(y_test, predictions)
classification_rep = classification_report(y_test, predictions)

# Print evaluation metrics
print("Decision Tree Model (Without Hyperparameter Tuning)")
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
# print("AUC-ROC Score:", auc_roc)
print("Confusion Matrix:\n", confusion)
print("Classification Report:\n", classification_rep)
```

Decision Tree Model (Without Hyperparameter Tuning)

Accuracy: 0.2833333333333333

Precision: 0.28688289827191615

Recall: 0.2833333333333333

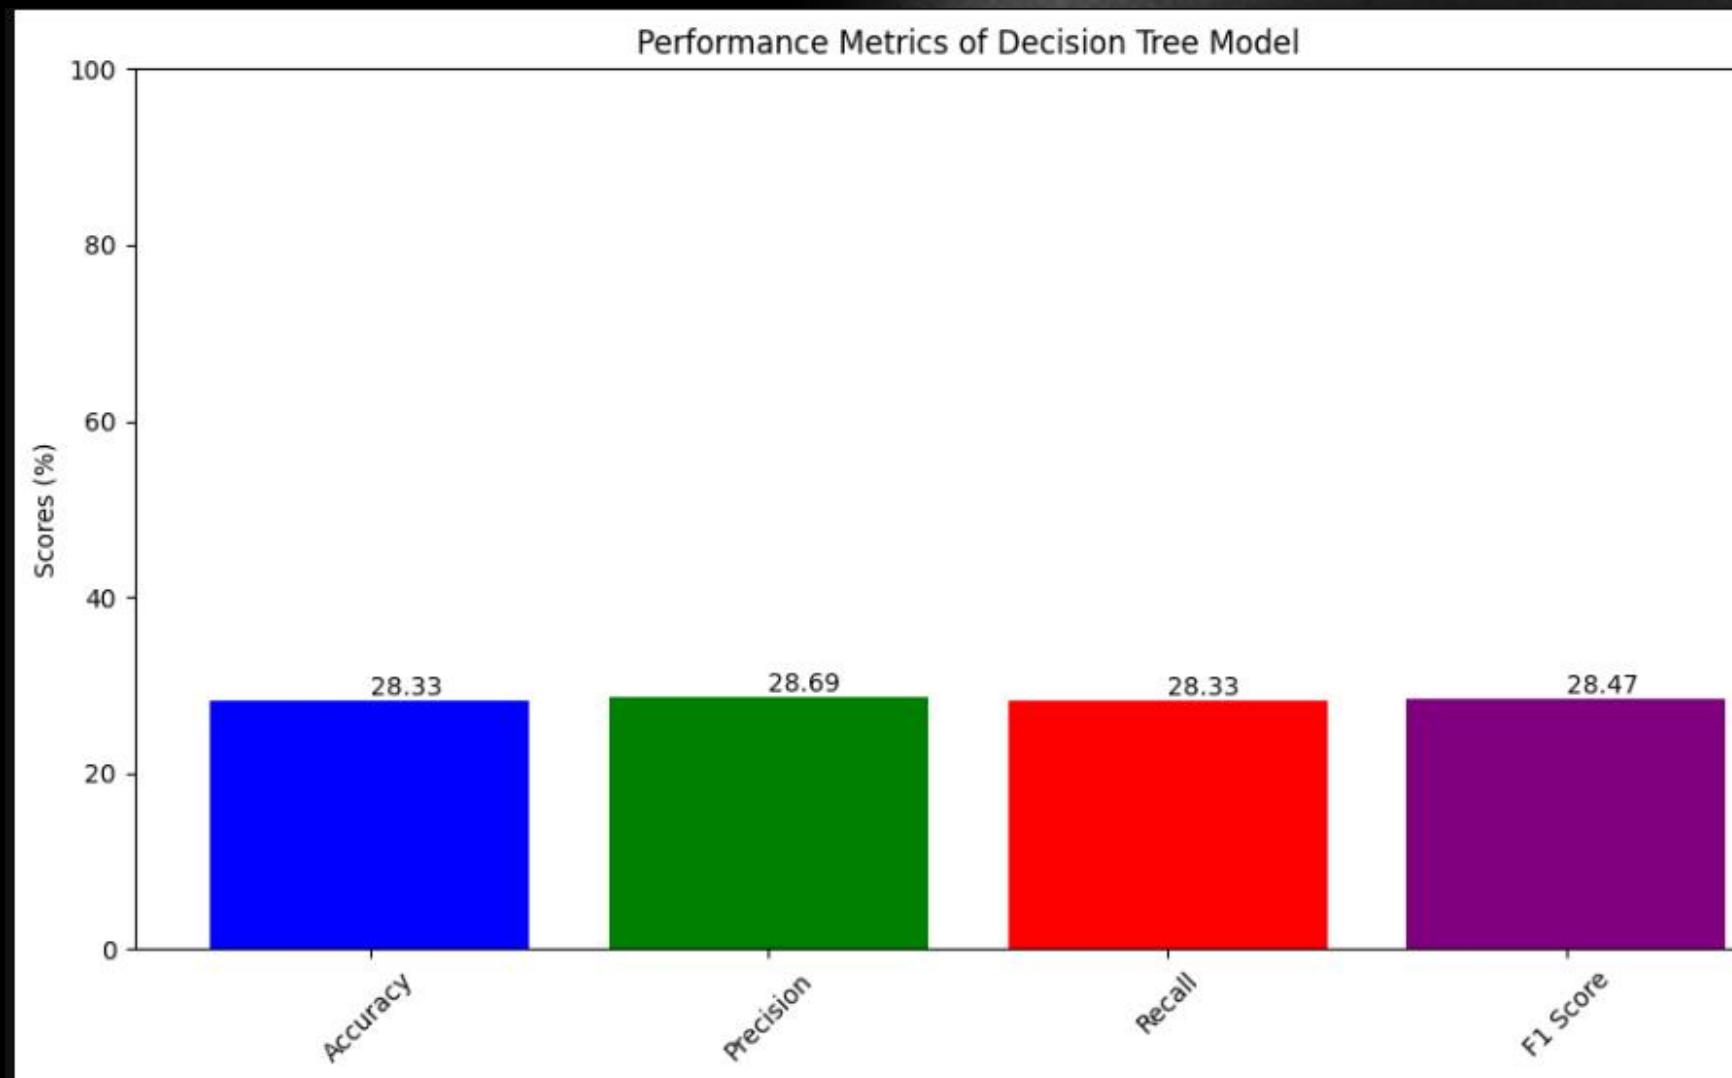
F1 Score: 0.28466715057830844

Confusion Matrix:

| |
|--------------------------------|
| [[0 0 0 0 0 0 1 0 0 0 0]] |
| [0 0 0 0 0 0 0 0 0 0 0] |
| [0 0 0 0 3 0 3 1 2 0 0] |
| [1 0 0 0 1 2 0 0 1 0 0] |
| [0 1 3 0 20 9 20 1 5 0 0] |
| [1 0 0 0 3 12 20 10 7 0 0] |
| [1 1 3 3 21 17 56 24 22 1 1] |
| [0 0 0 2 7 16 22 21 11 0 1] |
| [0 0 0 1 4 4 18 11 9 2 1] |
| [0 0 0 0 1 0 3 2 2 0 0] |
| [0 0 0 0 0 0 1 3 0 0 1]] |

Decision tree

Decision Tree:



Confusion Matrix for Decision Tree Model

A confusion matrix titled "Confusion Matrix for Decision Tree Model". The y-axis is labeled "True Labels" and the x-axis is labeled "Predicted Labels", both ranging from 0 to 10. The matrix shows the count of samples for each true label (0-10) across all predicted labels (0-10). The diagonal elements represent correct predictions, with the highest value being 56 for Predicted Label 6. A color scale on the right indicates the count, ranging from 0 (lightest blue) to 50 (darkest blue).

| Predicted Labels | | | | | | | | | | | |
|------------------|---|---|---|---|----|----|----|----|----|---|----|
| True Labels | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 2 | 0 | 0 | 0 | 0 | 0 | 3 | 0 | 3 | 1 | 2 | 0 |
| 3 | 1 | 0 | 0 | 0 | 0 | 1 | 2 | 0 | 0 | 1 | 0 |
| 4 | 0 | 1 | 3 | 0 | 0 | 20 | 9 | 20 | 1 | 5 | 0 |
| 5 | 1 | 0 | 0 | 0 | 0 | 3 | 12 | 20 | 10 | 7 | 0 |
| 6 | 1 | 1 | 3 | 3 | 21 | 17 | 56 | 24 | 22 | 1 | 1 |
| 7 | 0 | 0 | 0 | 2 | 7 | 16 | 22 | 21 | 11 | 0 | 1 |
| 8 | 0 | 0 | 0 | 1 | 4 | 4 | 18 | 11 | 9 | 2 | 1 |
| 9 | 0 | 0 | 0 | 0 | 1 | 0 | 3 | 2 | 2 | 0 | 0 |
| 10 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 3 | 0 | 0 |

Decision tree

Decision Tree:

we divided the dataset in 80 for training and testing.

tfidfVectorizer from sklearn.feature_extraction.text is used to convert the textdata into numericalfeatures suitable for modeltraining.

A DecisionTree model is instantiated and trained on the vectorizedtraining data (X_train_tfidf).
the trainedmodel is used to predict the scores of the vectorizedtestdata (X_test_tfidf).

Accuracy: The proportion of totalpredictions that were correct.

Precision (Weighed): The ability of the classifiernot to label a negativesample as positive, with consideration for theimbalance in classdistribution.

Recall (Weighed): The ability of the classifier to find all the positivesamples, again weighed for class imbalance.

F1 Score (Weighed): A weighed average of precision andrecall, which balance both metrics.

A confusion matrix is generated, showingtrue positives, falsepositives, truenegatives, and falsenegatives.

Neural Network

A neural network model is a computational framework inspired by the biological neural networks that constitute animal brains. It is capable of learning from observational data and has become a cornerstone for many applications in artificial intelligence, including image and speech recognition, natural language processing, and robotic control.

- Basic Structure:
- Neurons: The fundamental units of a neural network are artificial neurons or nodes. Each neuron receives input, processes it, and generates an output.
- Layers: Neurons are organized layers. There are three types of layers in a typical neural network:
- Input Layer: Receives the initial data.
- Hidden Layers: Intermediate layers that process inputs received from previous layers using weights, biases, and activation functions.
- Output Layer: Produces the final output of the model.

Neural Network

```
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
X_train, X_test, y_train, y_test = train_test_split(df['summary'], df['final_score'].astype(int), test_size=0.2, shuffle=True, random_state=26)
# Example DataFrame Loading

tfidf = TfidfVectorizer(max_features=10000)
X_train_tfidf = tfidf.fit_transform(X_train)
X_test_tfidf = tfidf.transform(X_test)

import tensorflow as tf

# Define the model architecture
model = tf.keras.Sequential([
    tf.keras.layers.Dense(512, activation='relu', input_shape=(X_train_tfidf.shape[1],)),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(256, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(11, activation='softmax') # Output layer for 11 classes
])

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
model.summary()

Model: "sequential_2"

```

| Layer (type) | Output Shape | Param # |
|---------------------|--------------|-----------|
| dense_6 (Dense) | (None, 512) | 5,120,512 |
| dropout_4 (Dropout) | (None, 512) | 0 |
| dense_7 (Dense) | (None, 256) | 131,328 |
| dropout_5 (Dropout) | (None, 256) | 0 |
| dense_8 (Dense) | (None, 11) | 2,827 |

```
Total params: 5,254,667 (20.04 MB)
Trainable params: 5,254,667 (20.04 MB)
Non-trainable params: 0 (0.00 B)
```

```
history = model.fit(X_train_tfidf, y_train, validation_split=0.1, epochs=10, batch_size=64)

Epoch 1/10
24/24 3s 46ms/step - accuracy: 0.2418 - loss: 2.2163 - val_accuracy: 0.3095 - val_loss: 1.7154
Epoch 2/10
24/24 1s 31ms/step - accuracy: 0.3315 - loss: 1.6915 - val_accuracy: 0.3274 - val_loss: 1.5997
Epoch 3/10
24/24 1s 31ms/step - accuracy: 0.4395 - loss: 1.4601 - val_accuracy: 0.3095 - val_loss: 1.5707
Epoch 4/10
24/24 1s 32ms/step - accuracy: 0.5683 - loss: 1.2169 - val_accuracy: 0.4048 - val_loss: 1.5079
Epoch 5/10
24/24 1s 30ms/step - accuracy: 0.7308 - loss: 0.8991 - val_accuracy: 0.4167 - val_loss: 1.5037
Epoch 6/10
24/24 1s 31ms/step - accuracy: 0.8646 - loss: 0.5545 - val_accuracy: 0.4167 - val_loss: 1.5814
Epoch 7/10
24/24 1s 31ms/step - accuracy: 0.9253 - loss: 0.3370 - val_accuracy: 0.4048 - val_loss: 1.7179
Epoch 8/10
24/24 1s 32ms/step - accuracy: 0.9493 - loss: 0.2261 - val_accuracy: 0.4345 - val_loss: 1.8432
Epoch 9/10
24/24 1s 33ms/step - accuracy: 0.9751 - loss: 0.1378 - val_accuracy: 0.3929 - val_loss: 1.9640
Epoch 10/10
24/24 1s 32ms/step - accuracy: 0.9863 - loss: 0.0898 - val_accuracy: 0.3929 - val_loss: 2.0674

loss, accuracy = model.evaluate(X_test_tfidf, y_test)
print(f"Test Accuracy: {accuracy*100:.2f}%")

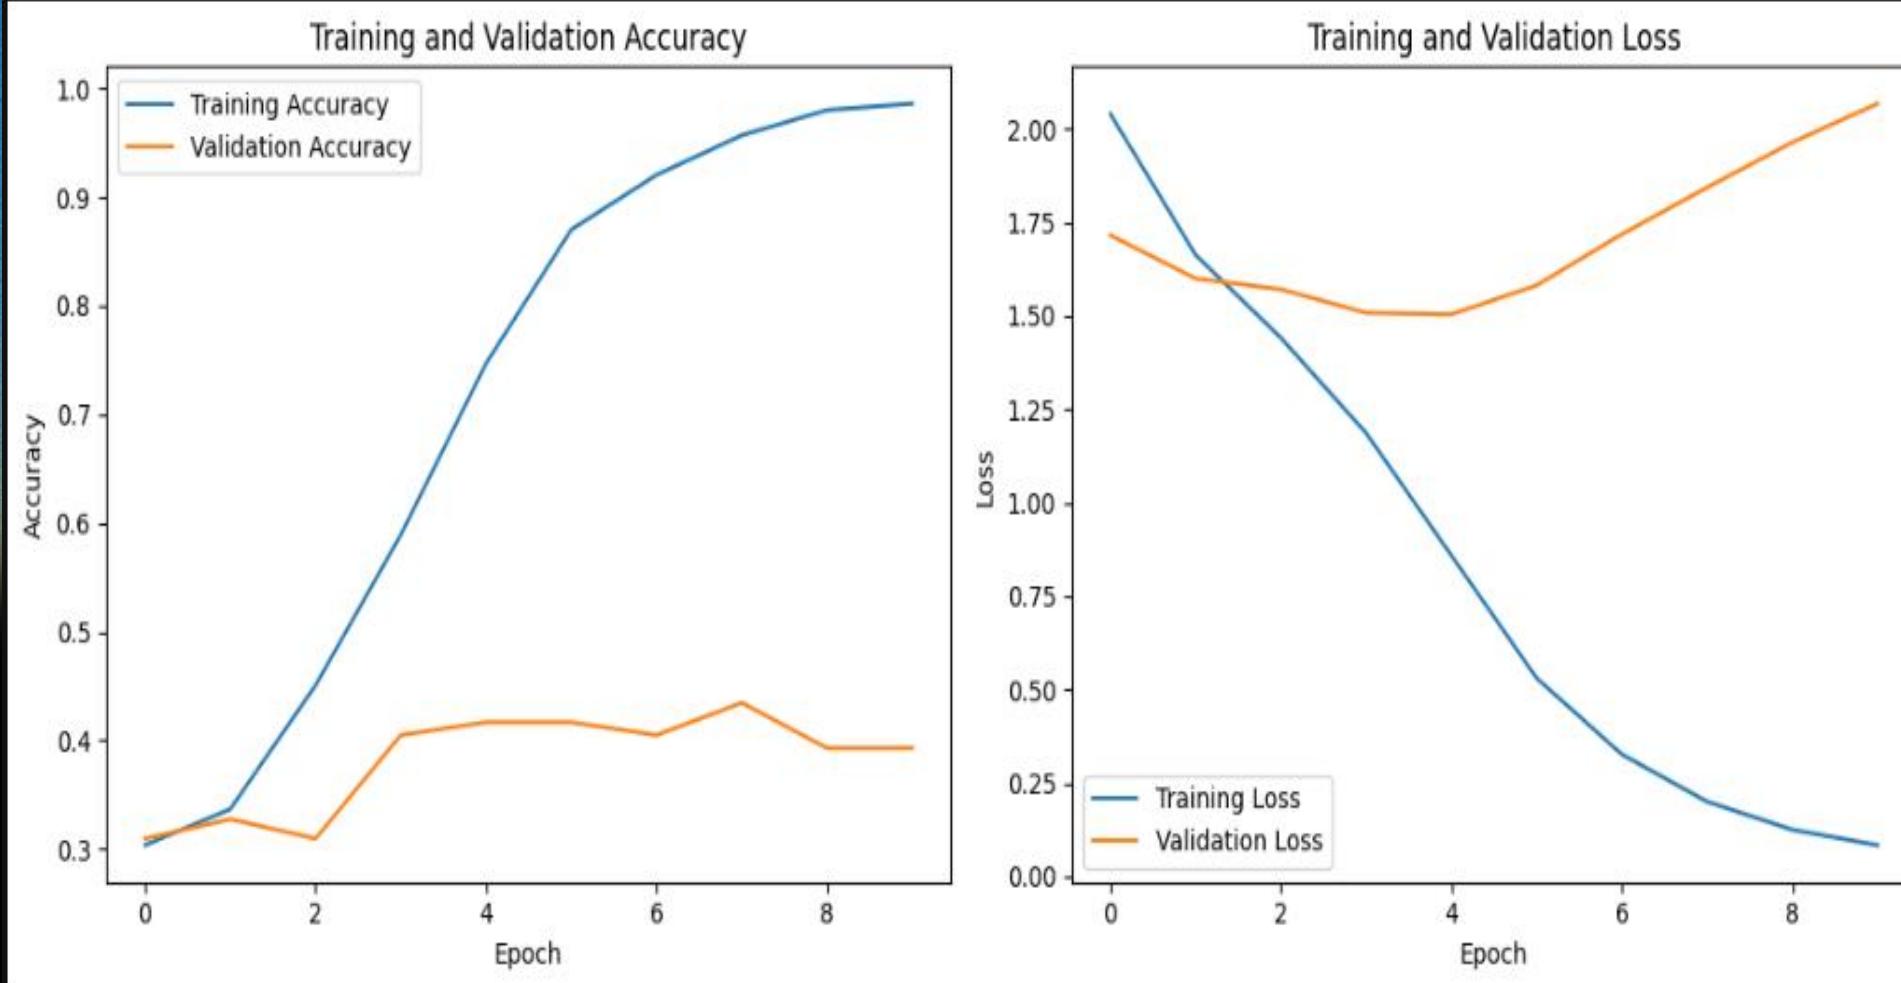
14/14 0s 2ms/step - accuracy: 0.3723 - loss: 2.0819
Test Accuracy: 35.71%
```

```
import matplotlib.pyplot as plt

# Ensure the figure is Large enough to hold both plots with clear visibility
plt.figure(figsize=(12, 5))

# Plot training and validation accuracy
plt.subplot(1, 2, 1) # 1 row, 2 columns, 1st subplot
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
```

Neural Network



```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
import numpy as np

# Predicting the probabilities
y_pred_probs = model.predict(X_test_tfidf)
# Converting probabilities to class labels
y_pred = np.argmax(y_pred_probs, axis=1)

14/14 ━━━━━━━━ 0s 7ms/step

# Calculate metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1 Score: {f1:.2f}")

Accuracy: 0.36
Precision: 0.33
Recall: 0.36
F1 Score: 0.34
```

After hyperparameter tuning: neural network

```
: import tensorflow as tf
from kerastuner.tuners import RandomSearch

def build_model(hp):
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Dense(
        units=hp.Int('units1', min_value=256, max_value=512, step=128),
        activation='relu',
        input_shape=(X_train_tfidf.shape[1],)
    ))
    model.add(tf.keras.layers.Dropout(rate=hp.Float('dropout1', min_value=0.4, max_value=0.6, step=0.1)))
    model.add(tf.keras.layers.Dense(
        units=hp.Int('units2', min_value=128, max_value=256, step=64),
        activation='relu'
    ))
    model.add(tf.keras.layers.Dropout(rate=hp.Float('dropout2', min_value=0.4, max_value=0.6, step=0.1)))
    model.add(tf.keras.layers.Dense(11, activation='softmax'))

    hp_learning_rate = hp.Float('learning_rate', min_value=1e-4, max_value=1e-2, sampling='LOG')
    model.compile(
        optimizer=tf.keras.optimizers.Adam(learning_rate=hp_learning_rate),
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy']
    )
    return model

:t tuner = RandomSearch(
    build_model,
    objective='val_accuracy',
    max_trials=10,
    executions_per_trial=1,
    directory='my_dir',
    project_name='keras_tuner_demo'
)

# Start hyperparameter search
tuner.search(X_train_tfidf, y_train, epochs=5, validation_split=0.1)

# Get the best hyperparameters
best_hps = tuner.get_best_hyperparameters(num_trials=1)[0]
print(f"""

The best number of units in the first densely-connected layer is {best_hps.get('units1')} and the best dropout rate is {best_hps.get('dropout1')}.

```

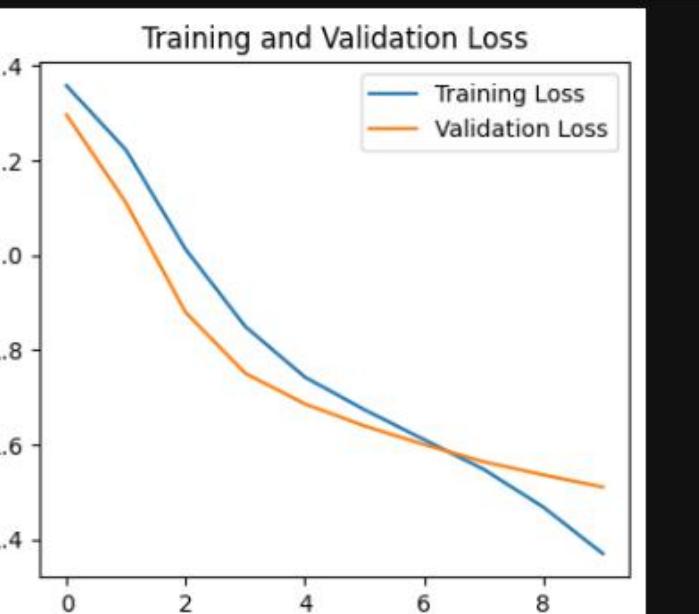
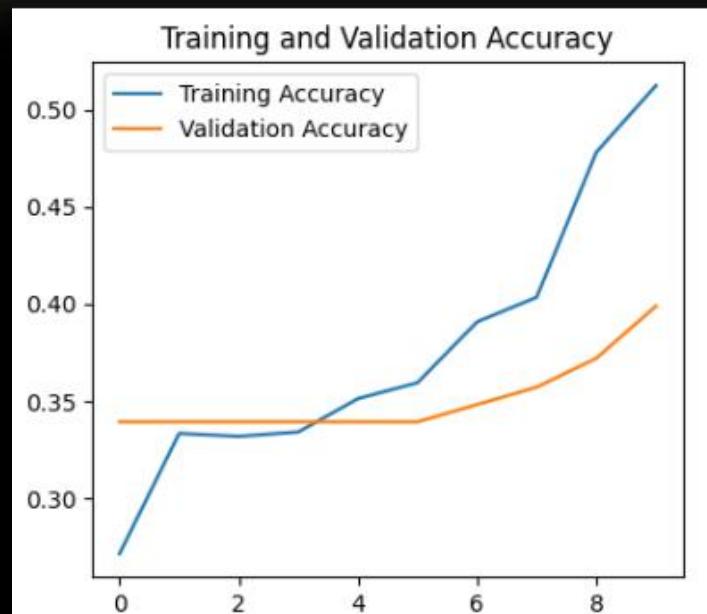
After hyperparameter tuning: neural network

```
import matplotlib.pyplot as plt

# Plot accuracy
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

# Plot loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.legend()

plt.show()
```



```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
import numpy as np
```

```
# Predicting the probabilities
y_pred_probs = model.predict(X_test_tfidf)
# Converting probabilities to class Labels
y_pred = np.argmax(y_pred_probs, axis=1)
```

```
14/14 ━━━━━━━━ 0s 7ms/step
```

```
# Calculate metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

print("Accuracy: {:.2f}")
print("Precision: {:.2f}")
print("Recall: {:.2f}")
print("F1 Score: {:.2f}")
```

```
Accuracy: 0.37
Precision: 0.28
Recall: 0.37
F1 Score: 0.29
```

After hyperparameter tuning: neural network

- Dense Layers: Fully connected layers with a variable number of units (`units1` and `units2`), both using ReLU activation functions. The number of units for these layers is chosen based on the hyperparameter search space defined for `units1` and `units2`.
- Dropout Layers: Used to prevent overfitting by randomly setting a fraction of input units to 0 during training. Dropout rates (`dropout1` and `dropout2`) are also variable.
- Output Layer: A softmax layer suitable for multi-class classification with 11 output classes.
- Optimization and Compilation:
- The learning rate for the Adam optimizer is varied within a specified range using a logarithmic scale (`sampling='LOG'`), which is effective for exploring orders of magnitude.
- The model uses `sparse_categorical_crossentropy` as the loss function, which is typical for multi-class classification tasks where labels are provided as integers.
- Kera's Tuner Configuration:
- Random Search: A tuner that performs a random search on hyperparameters, optimizing for `val_accuracy`.
- The search evaluates 10 different configurations (`max_trials=10`) with each configuration run once (`executions_per_trial=1`).
- The directory and project name specify where trial data is stored

After hyperparameter tuning: Logistic Regression

logistic regression

```
from sklearn.model_selection import GridSearchCV

# Define the parameter grid
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'solver': ['newton-cg', 'lbfgs', 'liblinear'],
    'penalty': ['l2']
}

# Ensure the model uses multi_class if it's a multi-class classification
logreg = LogisticRegression(max_iter=1000, multi_class='auto')

# Setup the grid search
grid_search = GridSearchCV(estimator=logreg, param_grid=param_grid, cv=5, verbose=1, scoring='accuracy', n_jobs=-1)

# Fit GridSearchCV
grid_search.fit(X_train_tfidf, y_train)

Fitting 5 folds for each of 18 candidates, totalling 90 fits
+   GridSearchCV      ⓘ ⓘ
  + estimator: LogisticRegression
    + LogisticRegression ⓘ
      | 

# Best parameters
best_params = grid_search.best_params_
print("Best parameters found:", best_params)

# Best model
best_model = grid_search.best_estimator_

Best parameters found: {'C': 1, 'penalty': 'l2', 'solver': 'newton-cg'}
```

```
] # Predictions with the best model
best_predictions = best_model.predict(X_test_tfidf)

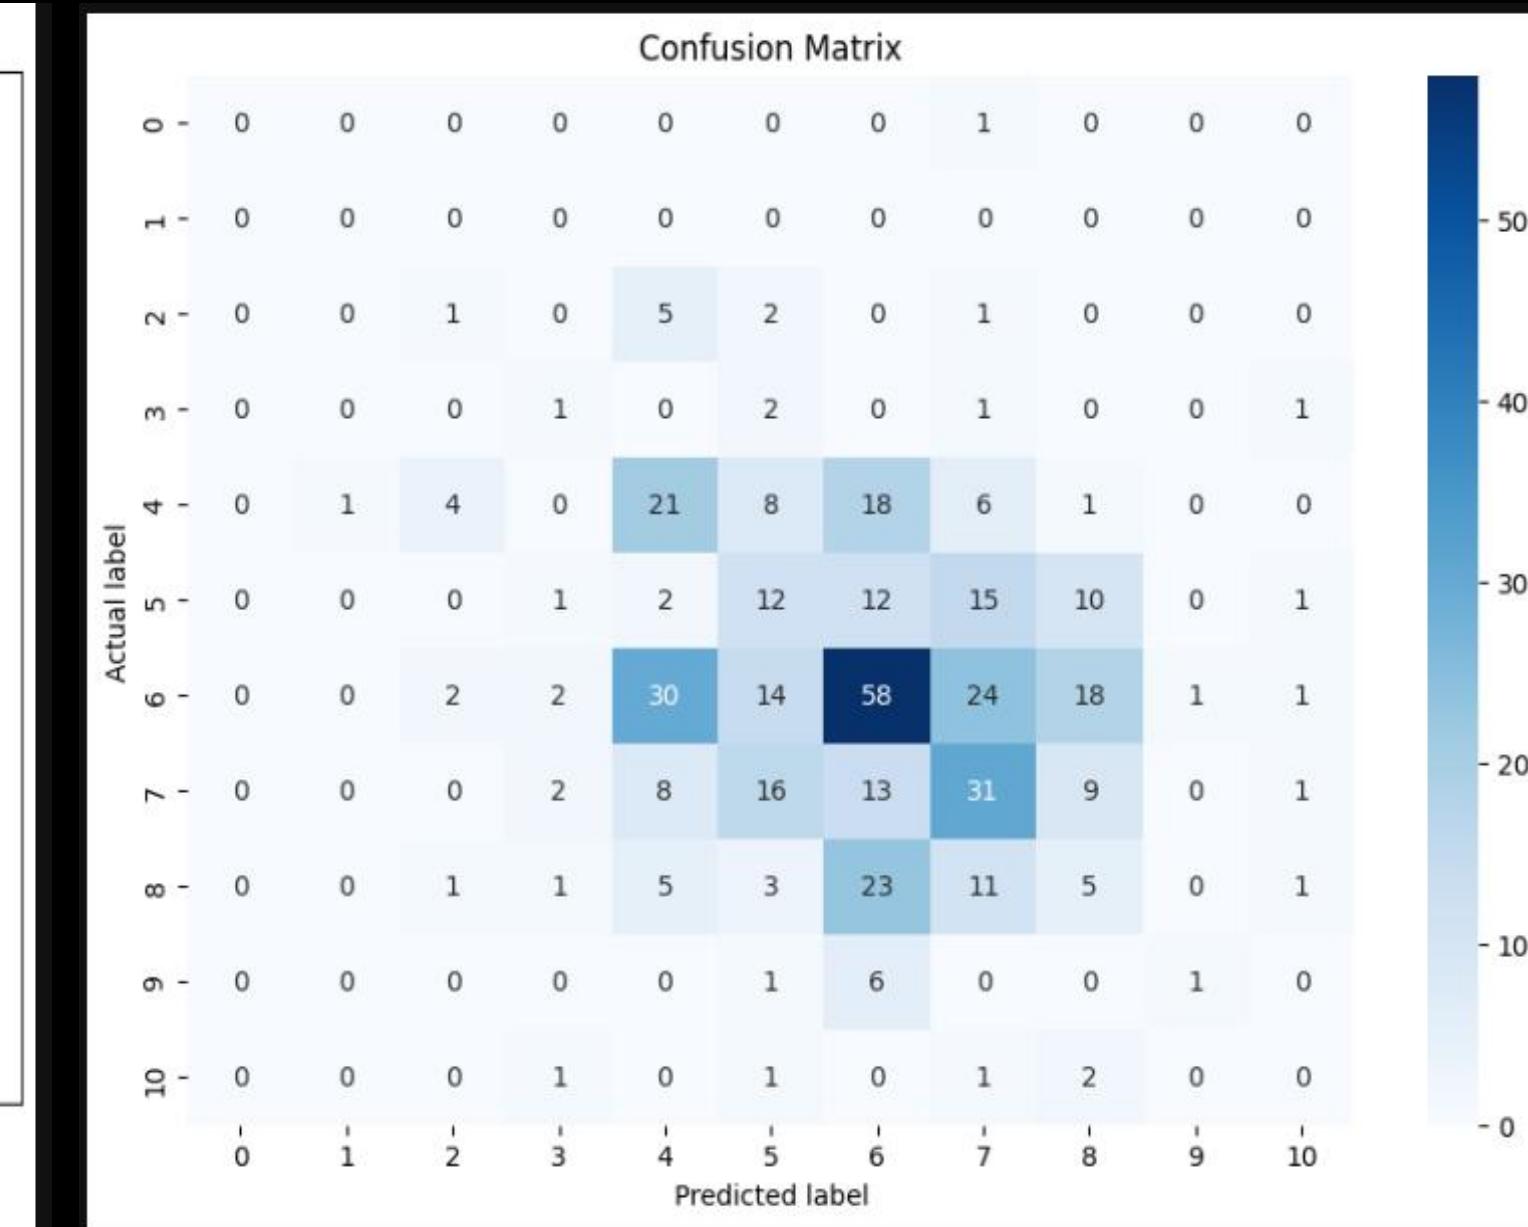
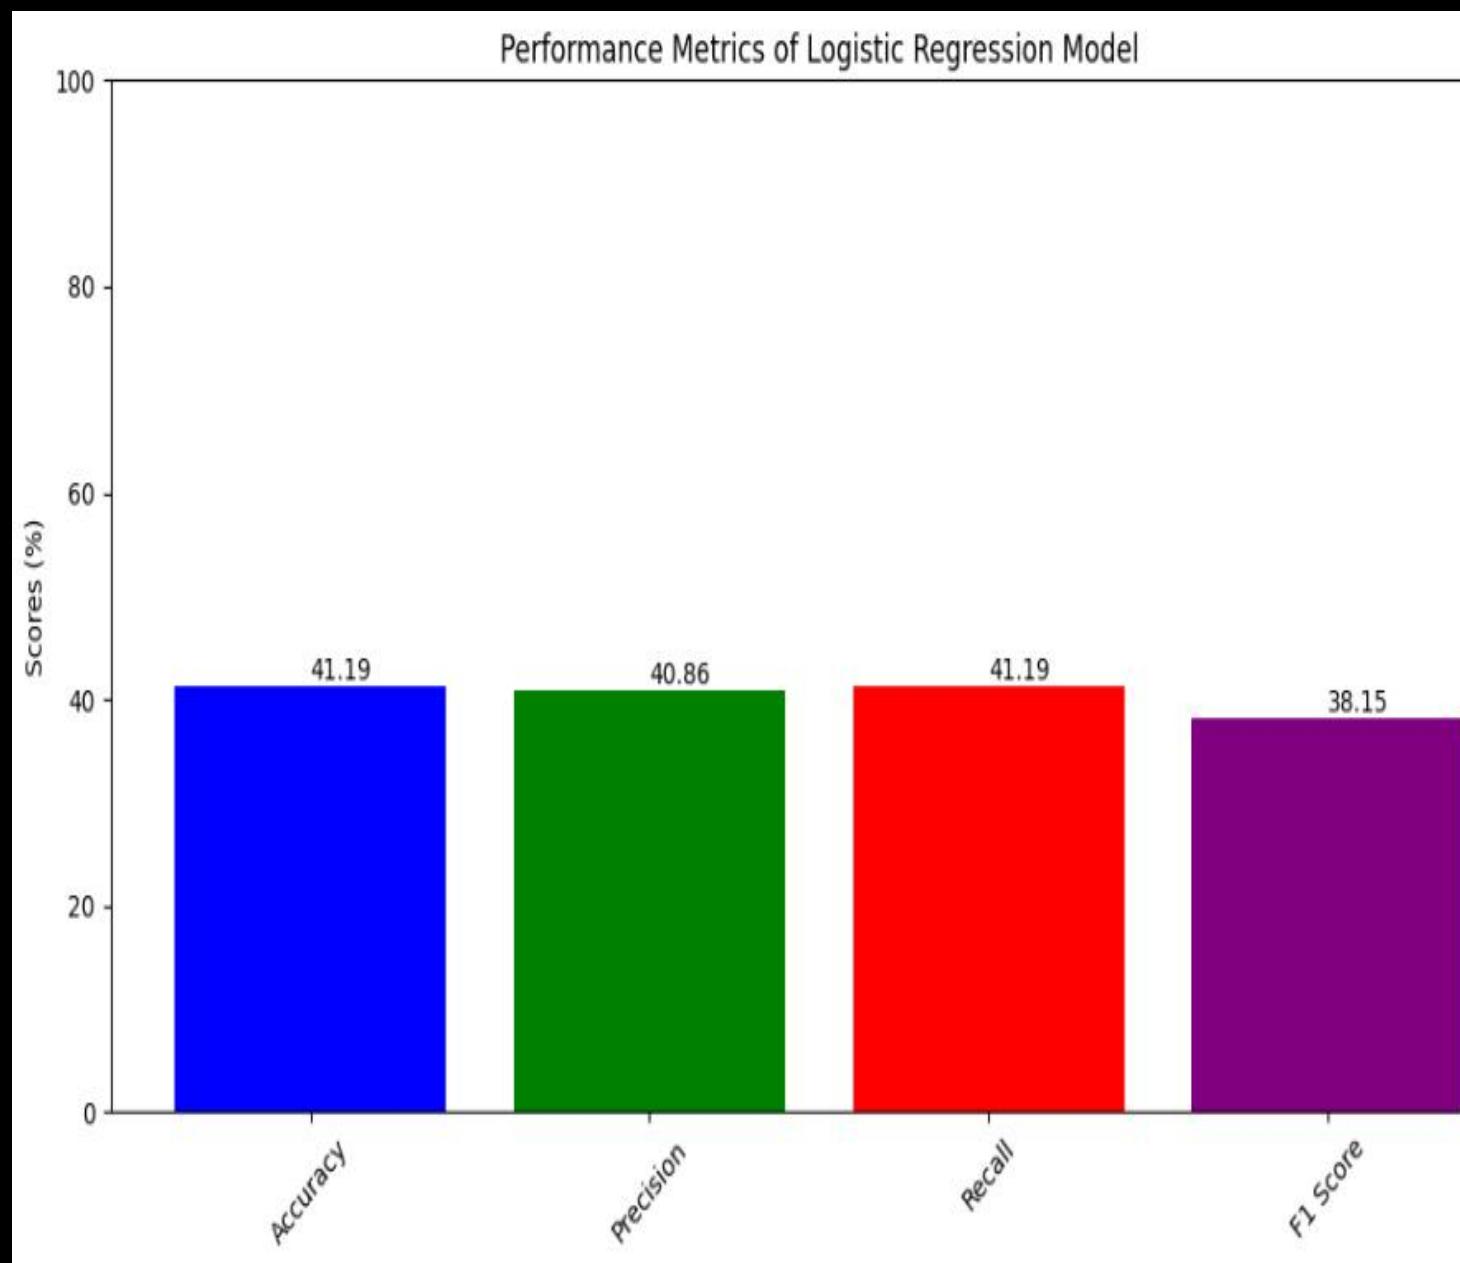
# Evaluation metrics
best_accuracy = accuracy_score(y_test, best_predictions)
best_precision = precision_score(y_test, best_predictions, average='weighted')
best_recall = recall_score(y_test, best_predictions, average='weighted')
best_f1 = f1_score(y_test, best_predictions, average='weighted')
best_confusion = confusion_matrix(y_test, best_predictions)
best_classification_rep = classification_report(y_test, best_predictions)

# Print evaluation metrics for the best model
print("Logistic Regression Model (With Hyperparameter Tuning)")
print("Accuracy:", best_accuracy)
print("Precision:", best_precision)
print("Recall:", best_recall)
print("F1 Score:", best_f1)
print("Confusion Matrix:\n", best_confusion)
print("Classification Report:\n", best_classification_rep)
```

```
Logistic Regression Model (With Hyperparameter Tuning)
Accuracy: 0.40714285714285714
Precision: 0.3922068647678404
Recall: 0.40714285714285714
F1 Score: 0.36319090053175274
Confusion Matrix:
[[ 0  0  0  1  0  0  0  0  0  0]
 [ 0  0  0  3  2  4  0  0  0  0]
 [ 0  0  0  0  3  1  1  0  0  0]
 [ 0  0  0  21 12 26  0  0  0  0]
 [ 0  0  0  0 12 31 10  0  0  0]
 [ 0  0  0  15  6 95 33  1  0  0]
 [ 0  0  0  1  2 35 41  1  0  0]
 [ 0  0  0  0 27 21  2  0  0  0]
 [ 0  0  0  0  7  1  0  0  0  0]
 [ 0  0  0  0  0  2  3  0  0  0]]
```

| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
|--|-----------|--------|----------|---------|

After hyperparameter tuning: Logistic Regression



After hyperparameter tuning: Logistic Regression

Setting up Gridsearch CV:

- ParameterGrid: A dictionary (`param_grid`) defines the space over which parameters should be searched. This includes:
 - C: Regularization strength, which inversely scales with overfitting. Several values are tested from very weak to very strong regularization.
 - Solver: Different algorithms for solving the optimization problem (`newton-cg`, `lbfgs`, `liblinear`). Each solver works differently under various conditions and is suitable for different types of data.
 - Penalty: The type of regularization applied (`l2` indicates L2 regularization).
- Logistic Regression Configuration: The base model is configured with a high `max_iter` to ensure convergence and set to handle multi-class classification automatically (`multi_class='auto'`).
- Grid Search Configuration: `GridSearchCV` is initialized with the `LogisticRegression` estimator, the parameter grid, and the strategy for cross-validation (`cv=5` indicates 5-fold cross-validation). It aims to maximize accuracy (`scoring='accuracy'`), and it is set to perform computations in parallel (`n_jobs=-1` uses all available CPU cores).

After hyperparameter tuning: Decision Tree

decision tree

```
from sklearn.model_selection import GridSearchCV

# Define the parameter grid
param_grid = {
    'max_depth': [None, 10, 20, 30, 50, 100],
    'min_samples_split': [2, 5, 10, 20],
    'min_samples_leaf': [1, 2, 5, 10]
}

# Initialize the GridSearchCV object
grid_search = GridSearchCV(estimator=DecisionTreeClassifier(), param_grid=param_grid, cv=5, scoring='accuracy', verbose=1, n_jobs=-1)

# Fit GridSearchCV
grid_search.fit(X_train_tfidf, y_train)

Fitting 5 folds for each of 96 candidates, totalling 480 fits
>      GridSearchCV      ⓘ ⓘ
> estimator: DecisionTreeClassifier
  > DecisionTreeClassifier ⓘ
|
```



```
# Best parameters found
best_params = grid_search.best_params_
print("Best parameters found:", best_params)

# Best model
best_dt_model = grid_search.best_estimator_

Best parameters found: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 5}
```

```
# Best parameters found
best_params = grid_search.best_params_
print("Best parameters found:", best_params)

# Best model
best_dt_model = grid_search.best_estimator_

Best parameters found: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 5}

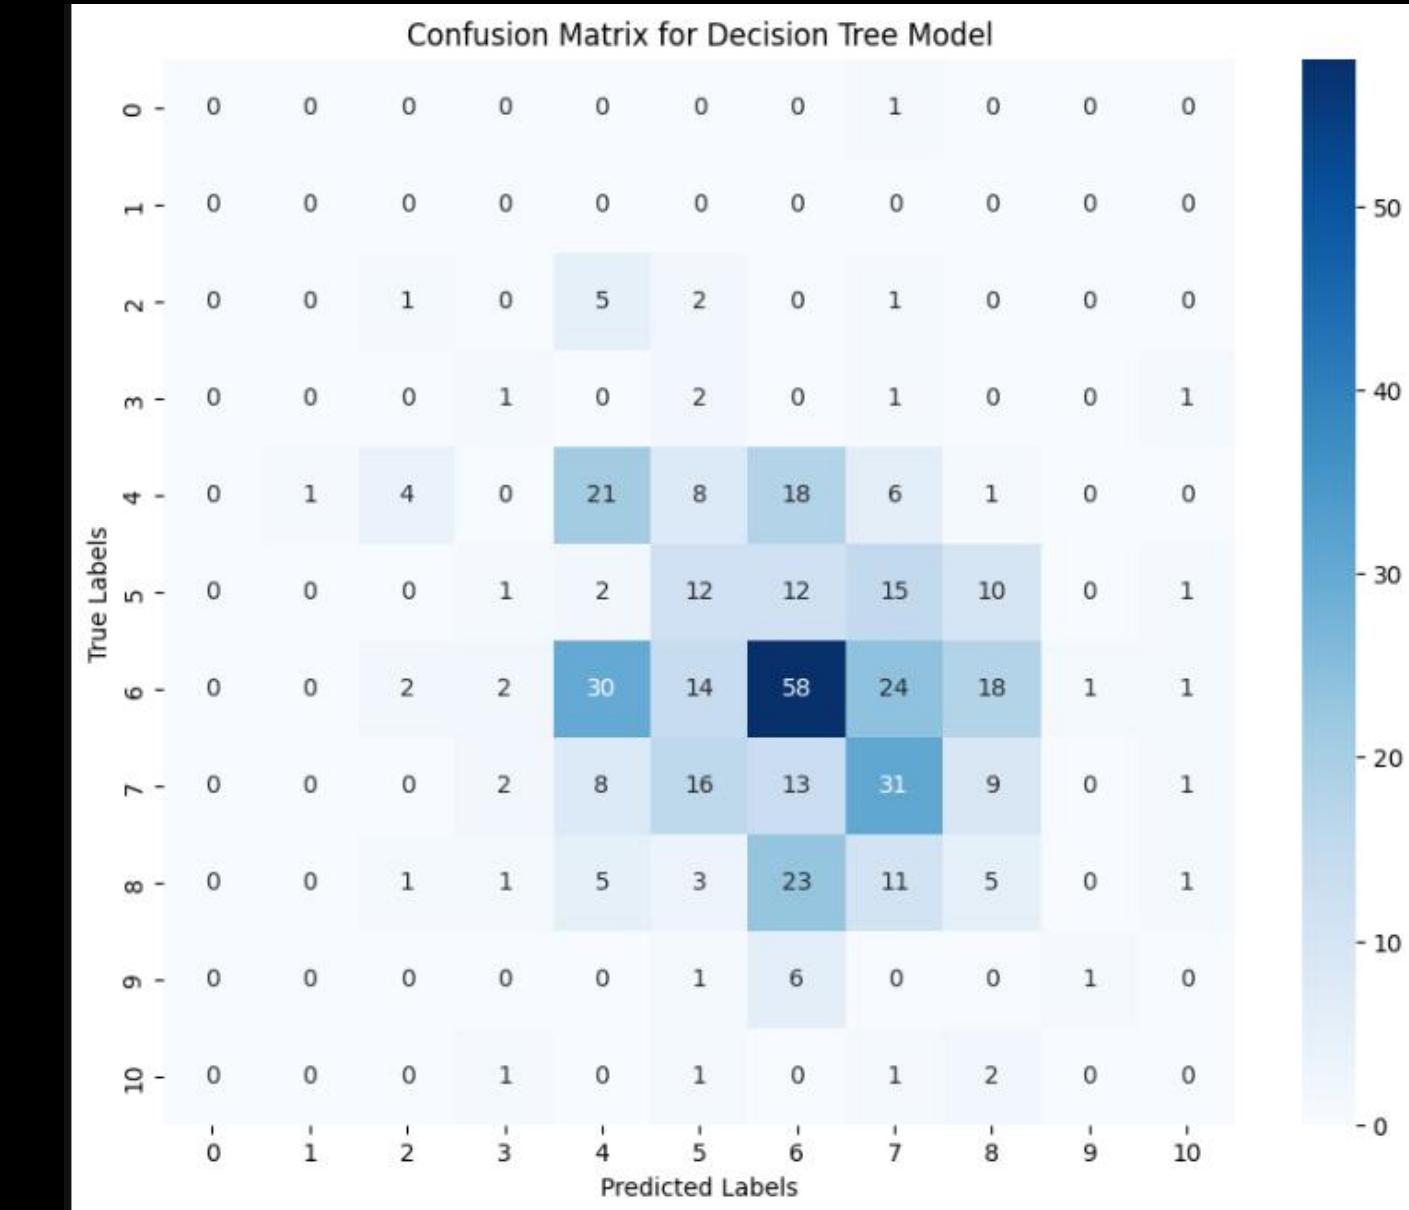
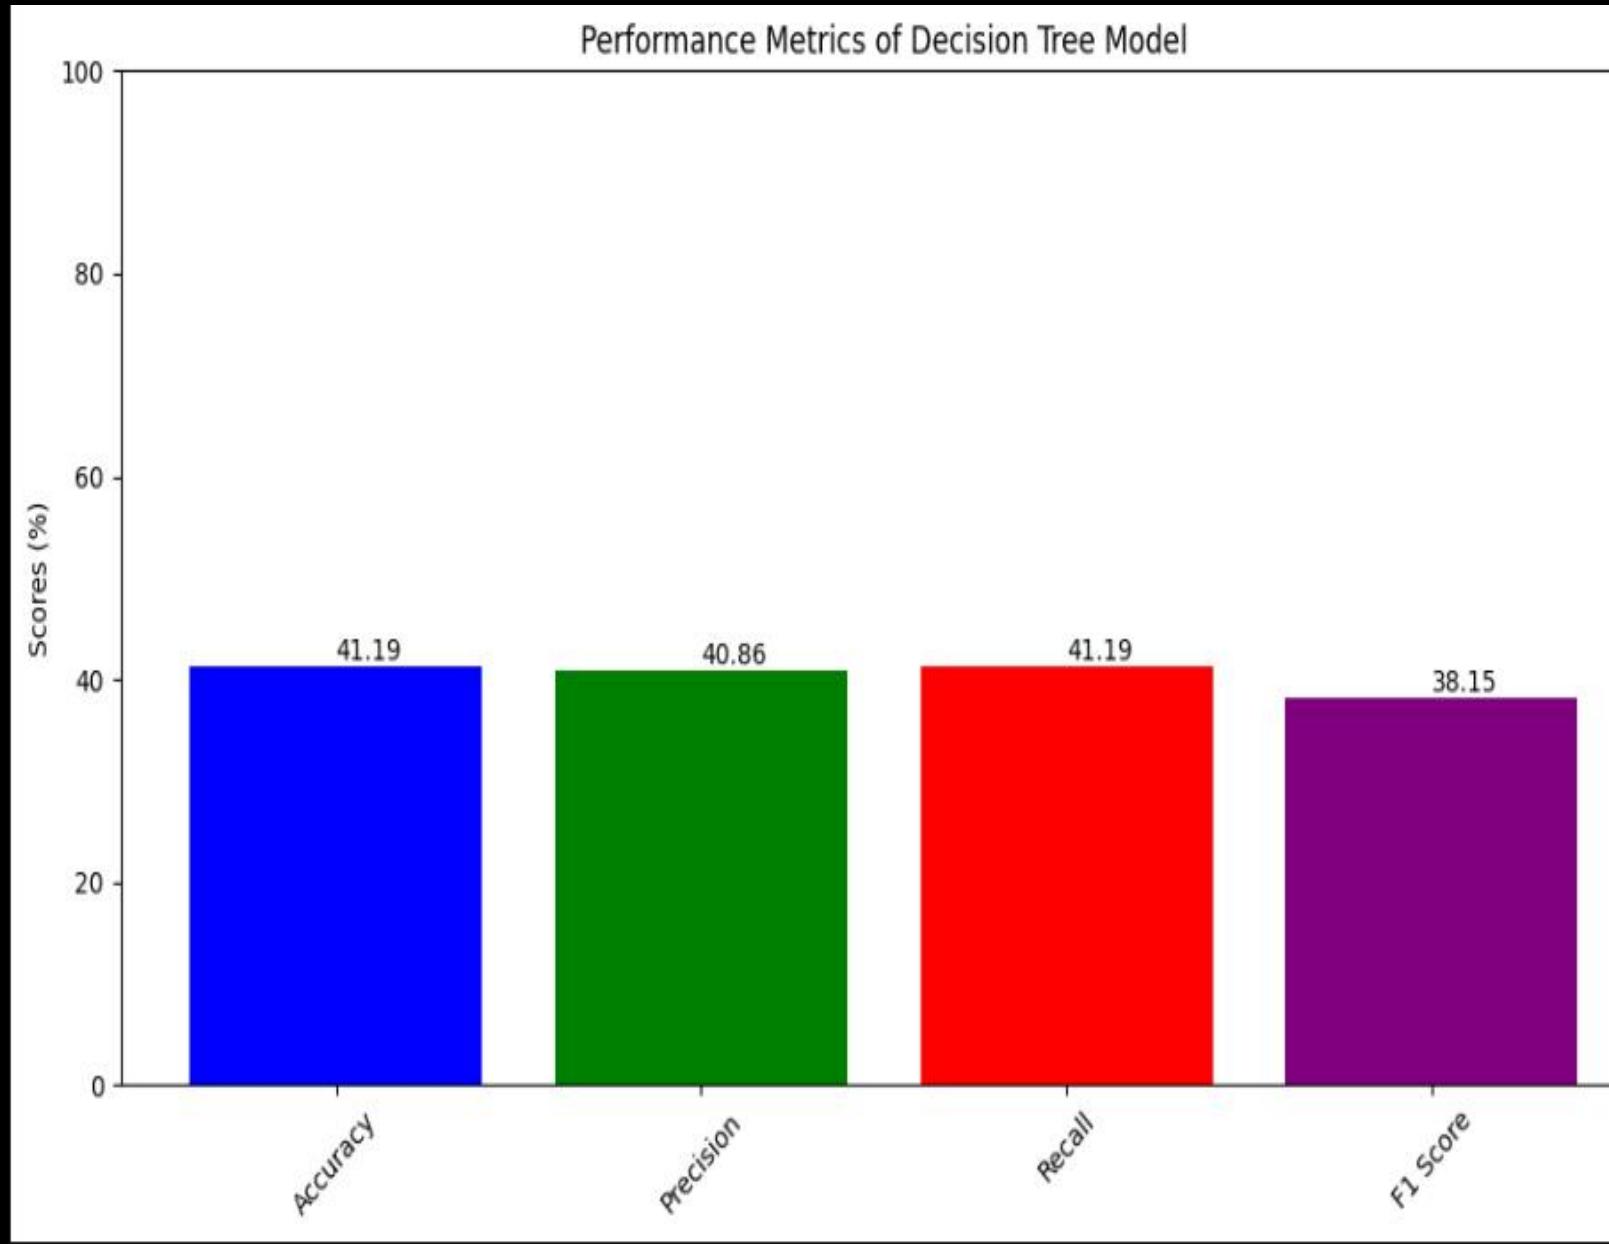
# Predictions with the best model
best_predictions = best_dt_model.predict(X_test_tfidf)

# Evaluation metrics
best_accuracy = accuracy_score(y_test, best_predictions)
best_precision = precision_score(y_test, best_predictions, average='weighted')
best_recall = recall_score(y_test, best_predictions, average='weighted')
best_f1 = f1_score(y_test, best_predictions, average='weighted')
best_confusion = confusion_matrix(y_test, best_predictions)
best_classification_rep = classification_report(y_test, best_predictions)

# Print evaluation metrics for the best model
print("Decision Tree Model (With Hyperparameter Tuning)")
print("Accuracy:", best_accuracy)
print("Precision:", best_precision)
print("Recall:", best_recall)
print("F1 Score:", best_f1)
print("Confusion Matrix:\n", best_confusion)
print("Classification Report:\n", best_classification_rep)

Decision Tree Model (With Hyperparameter Tuning)
Accuracy: 0.33095238095238094
Precision: 0.3157137307867451
Recall: 0.33095238095238094
F1 Score: 0.3106383630137287
Confusion Matrix:
[[ 0  0  0  0  0  0  1  0  0  0]
 [ 0  0  0  4  0  1  4  0  0  0]
 [ 0  0  0  0  1  1  3  0  0  0]
 [ 0  5  0  13  8  24  8  1  0  0]
 [ 0  0  0  2  10  15  26  0  0  0]
 [ 0  0  0  17  11  67  47  8  0  0]
 [ 0  0  0  2  7  27  42  2  0  0]
 [ 0  0  0  5  6  17  15  7  0  0]
 [ 0  0  0  2  0  4  0  2  0  0]
 [ 0  0  0  0  1  3  1  0  0]]
```

After hyperparameter tuning: Decision Tree



After hyperparameter tuning: Decision Tree

Setting up Gridsearch CV:

- Parameter Grid Definition: A dictionary (param_grid) is created to specify the range of values for various hyperparameters of the DecisionTreeClassifier:
 - max_depth: Controls the maximum depth (or levels) of the tree. Deep trees can model more complex patterns but are prone to overfitting.
 - min_samples_split: The minimum number of samples required to split an internal node. Higher values prevent the model from learning overly specific patterns, thus lowering the risk of overfitting.
 - min_samples_leaf: The minimum number of samples a leaf node must have. This parameter further helps in controlling overfitting.
- GridSearchCV Configuration: An instance of GridSearchCV is initialized with the Decision Tree Classifier as the estimator. It uses the defined parameter grid to explore different configurations. The search strategy includes:
 - cv=5: 5-fold cross-validation to ensure the model's performance is robust across different subsets of the dataset.
 - scoring='accuracy': Optimization criterion is the accuracy of the predictions.
 - verbose=1: Outputs detailed logging information during the training process.
 - n_jobs=-1: Utilizes all available CPU cores to parallelize the training.

Project Management

| Person | Task | Description | Contribution | Issues faced |
|-------------------|---------------------------------------|--|--------------|---|
| Sai Gottumukkala | Visualize and Pre-process the dataset | To clean the essay dataset using special symbols, slashes etc. | 100% | Data available was less, further after cleaning and summarization, more data was lost resulting in degradation of ML models from the summary. |
| Sai Charan Merugu | Model analysis using Neural Network | To design a neural network model that can classify the grade | 100% | Data available was less, further after cleaning and summarization, more data was lost resulting in degradation of ML models from the summary. |

| Person | Task | Description | Contribution | Issues faced |
|--------------------|--|---|--------------|---|
| Swetha Guntupalli | Model analysis using Logistic Regression | To design a model that can grade the input summarized essay text. | 100% | Data available was less, further after cleaning and summarization, more data was lost resulting in degradation of ML models from the summary. |
| Karishma Bollineni | Model analysis using decision tree | To train and implement decision tree, neural network and measure the metrics. | 100% | Data available was less, further after cleaning and summarization, more data was lost resulting in degradation of ML models from the summary. |

References

-  **The Hewlett Foundation: Automated Essay Scoring**
[“<https://www.kaggle.com/c/asap-aes/data>”](https://www.kaggle.com/c/asap-aes/data)
-  **Text Summarization**
[“<https://cs.nyu.edu/~kcho/DMQA/>”](https://cs.nyu.edu/~kcho/DMQA/)
-  **Keras**
[“\[https://keras.io/keras_tuner/\]\(https://keras.io/keras_tuner/\)”](https://keras.io/keras_tuner/)

Thank you